

# initial-results-and-the-code

March 30, 2024

This notebook presents the data cleaning, exploratory data analysis and feature engineering using data from [Gapminder](#) that will be used for linear, polynomial and logarithmic regression in another notebook.

## 1 Import packages and data

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns

# load the dataset (csv files)
data = pd.read_csv("https://raw.githubusercontent.com/BME1478H/Winter2022class/
↳master/data/world-data-gapminder.csv")
```

## 2 Description of data set and a summary of its attributes

This data set has been generated using data from [the Gapminder website](#), which focuses on gathering and sharing statistics and other information about social, economic and environmental development at local, national and global levels.

This particular data set describes the values of several parameters (see the list below) between 1800 and 2018 for a total of 192 countries. The parameters included in the data set and the column name of the dataframe are as follows:

- Country (country): Describes the country name
- Year (year): Describes the year to which the data belongs
- Population(population): Describes the number of population.
- Region (region): Describes the region to which the country belongs \*Sub Region (sub\_region): Describes the sub-region to which the country belongs.
- Income Group (income\_group): Classifies the people into income levels or classes.
- Life expectancy (life\_expectancy): Describes the life expectancy for a given country in a given year
- Income (income): Describes the Income per person in dollars for a given country in a given year
- Children per woman (children\_per-woman): Describes average number of children per woman
- Child Mortality (child-mortality): Describe number of child mortality in a country.
- Population Density (pop\_density): Describes the population density in a country.

- CO2 emissions per capita (co2\_per\_capita): Describes the CO2 emissions in tonnes per person for a given country in a given year
- Years in school for Men (years\_in\_school\_men): Describe the number of years on average a men spent in school
- Years in school for Women (years\_in\_school\_women): Describe the number of years on average an women spent in school

Let's have a look at the data.

Let's check the data types and the number of samples for each column:

[2]: `#Dataset chekup`  
`data.sample(n=10)`

```
[2]:      country  year  population  region \
17124  Jamaica  1842    400000  Americas
29762  Senegal  1997    920000   Africa
219    Albania  1800    410000  Europe
8220   Costa Rica 1917    400000  Americas
24398   Nepal  1889    5240000   Asia
34974   Togo   1953    1450000   Africa
1961    Austria 2009    8370000   Europe
1638   Australia 1905    4020000  Oceania
14796   Haiti  1923    2230000  Americas
27204  Paraguay 1848     318000  Americas

      sub_region  income_group  life_expectancy  income \
17124 Latin America and the Caribbean  Upper middle      34.2      977
29762          Sub-Saharan Africa      Low      58.8     1760
219          Southern Europe  Upper middle      35.4      667
8220 Latin America and the Caribbean  Upper middle      35.6     2300
24398          Southern Asia      Low      33.8      758
34974          Sub-Saharan Africa      Low      39.0     1080
1961          Western Europe      High      80.3    42500
1638 Australia and New Zealand      High      52.7     7140
14796 Latin America and the Caribbean      Low      28.5     1810
27204 Latin America and the Caribbean  Upper middle      35.5     1190

      children_per_woman  child_mortality  pop_density  co2_per_capita \
17124                5.13            390.0         NaN         NaN
29762                5.77            143.0         47.8         0.3550
219                4.60            375.0         NaN         NaN
8220                6.71            366.0         NaN         NaN
24398                6.15            407.0         NaN         NaN
34974                6.33            308.0         26.7         0.0202
1961                1.39              4.5        102.0         7.4900
1638                3.51            117.0         NaN         3.0000
14796                6.31            462.0         NaN         NaN
```

|       |                     |                       |     |     |
|-------|---------------------|-----------------------|-----|-----|
| 27204 | 6.49                | 374.0                 | NaN | NaN |
|       | years_in_school_men | years_in_school_women |     |     |
| 17124 | NaN                 | NaN                   |     |     |
| 29762 | 3.34                | 1.82                  |     |     |
| 219   | NaN                 | NaN                   |     |     |
| 8220  | NaN                 | NaN                   |     |     |
| 24398 | NaN                 | NaN                   |     |     |
| 34974 | NaN                 | NaN                   |     |     |
| 1961  | 12.40               | 12.40                 |     |     |
| 1638  | NaN                 | NaN                   |     |     |
| 14796 | NaN                 | NaN                   |     |     |
| 27204 | NaN                 | NaN                   |     |     |

```
[3]: # Display the first few rows of the dataset to understand its structure
print(data.head())
```

|   | country     | year | population | region | sub_region    | income_group | \ |
|---|-------------|------|------------|--------|---------------|--------------|---|
| 0 | Afghanistan | 1800 | 3280000    | Asia   | Southern Asia | Low          |   |
| 1 | Afghanistan | 1801 | 3280000    | Asia   | Southern Asia | Low          |   |
| 2 | Afghanistan | 1802 | 3280000    | Asia   | Southern Asia | Low          |   |
| 3 | Afghanistan | 1803 | 3280000    | Asia   | Southern Asia | Low          |   |
| 4 | Afghanistan | 1804 | 3280000    | Asia   | Southern Asia | Low          |   |

|   | life_expectancy | income | children_per_woman | child_mortality | pop_density | \ |
|---|-----------------|--------|--------------------|-----------------|-------------|---|
| 0 | 28.2            | 603    | 7.0                | 469.0           | NaN         |   |
| 1 | 28.2            | 603    | 7.0                | 469.0           | NaN         |   |
| 2 | 28.2            | 603    | 7.0                | 469.0           | NaN         |   |
| 3 | 28.2            | 603    | 7.0                | 469.0           | NaN         |   |
| 4 | 28.2            | 603    | 7.0                | 469.0           | NaN         |   |

|   | co2_per_capita | years_in_school_men | years_in_school_women |
|---|----------------|---------------------|-----------------------|
| 0 | NaN            | NaN                 | NaN                   |
| 1 | NaN            | NaN                 | NaN                   |
| 2 | NaN            | NaN                 | NaN                   |
| 3 | NaN            | NaN                 | NaN                   |
| 4 | NaN            | NaN                 | NaN                   |

```
[4]: # Display data types
print(data.dtypes)
```

|                 |         |
|-----------------|---------|
| country         | object  |
| year            | int64   |
| population      | int64   |
| region          | object  |
| sub_region      | object  |
| income_group    | object  |
| life_expectancy | float64 |

```

income                int64
children_per_woman    float64
child_mortality       float64
pop_density           float64
co2_per_capita        float64
years_in_school_men   float64
years_in_school_women float64
dtype: object

```

```

[5]: # List of columns for reference in the analyses below
      print(data.columns)

```

```

Index(['country', 'year', 'population', 'region', 'sub_region', 'income_group',
      'life_expectancy', 'income', 'children_per_woman', 'child_mortality',
      'pop_density', 'co2_per_capita', 'years_in_school_men',
      'years_in_school_women'],
      dtype='object')

```

```

[6]: # For categorical attributes, let's look at the frequency of countries or other
      ↪ categorical fields
      print(data["country"].value_counts()) # Assuming 'country' is one of the
      ↪ columns

```

```

United States    438
Afghanistan      219
Panama           219
New Zealand      219
Nicaragua        219
...
Greece           219
Grenada          219
Guatemala        219
Guinea           219
Zimbabwe         219
Name: country, Length: 178, dtype: int64

```

```

[7]: # Basic information about the dataset
      print(data.info())
      #data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39201 entries, 0 to 39200
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country                39201 non-null  object
1   year                  39201 non-null  int64
2   population             39201 non-null  int64

```

|    |                       |                |         |
|----|-----------------------|----------------|---------|
| 3  | region                | 39201 non-null | object  |
| 4  | sub_region            | 39201 non-null | object  |
| 5  | income_group          | 39201 non-null | object  |
| 6  | life_expectancy       | 39201 non-null | float64 |
| 7  | income                | 39201 non-null | int64   |
| 8  | children_per_woman    | 39201 non-null | float64 |
| 9  | child_mortality       | 39199 non-null | float64 |
| 10 | pop_density           | 12351 non-null | float64 |
| 11 | co2_per_capita        | 16500 non-null | float64 |
| 12 | years_in_school_men   | 8234 non-null  | float64 |
| 13 | years_in_school_women | 8234 non-null  | float64 |

dtypes: float64(7), int64(3), object(4)

memory usage: 4.2+ MB

None

### 3 Initial plan for data exploration

The initial plan for data exploration is as follows:

\*Data cleaning: The data cleaning process is split into following two main actions . . . Null values: Understand the reason why there are null values to find out the best way to deal with them.

Outliers: Similarly to the null values, the first step is to understand the presence of outliers as well as to find out if removing them is a good idea or it's actually valuable data.

\*Feature extraction: Understand the relationship between the different features, perform transformation to help improve those relationships and perform Principal Component Analysis to understand how some of the features explain the HDI index variance

\*Feature extraction: Understand the relationship between the different features, perform transformation to help improve those relationships and perform Principal Component Analysis to understand how some of the features explain the HDI index variance

\*Hypothesis testing: Formulate three hypotheses and test one of them

Next, Let's find out how many null values there are:

```
[8]: # To find out the number of Null values
data.isnull().sum()
```

```
[8]: country          0
year                0
population          0
region              0
sub_region          0
income_group        0
life_expectancy     0
income              0
children_per_woman  0
child_mortality     2
```

```

pop_density          26850
co2_per_capita        22701
years_in_school_men   30967
years_in_school_women 30967
dtype: int64

```

## 4 Data cleaning

### 4.1 Null values

Data cleaning is always tricky and this dataset isn't an exception. In fact, it's even trickier because the number of samples, 39202, is spread accross 218 different years. This means that there are 218 different sub-datasets with around 205 elements each. Therefore, the sample size is relatively small and the impact of each value in the final metrics is greater. One needs to be careful and undrestand the reason behind the NaN values before replacing or deleting them.

Prior to the analysis, I considered the following techniques to deal with null values:

- Remove them, provided that we don't lose a lot of data
- Impute data using the one of the following options:
  - Mean value
  - Most common value
  - Interpolation provided that there is enough data to make a prediction

However, let's starts by checking which are the null values and try to understand why there are missing values.

```
[9]: data[data.isna().any(axis=1)]
```

```

[9]:
   country  year  population  region  sub_region income_group \
0  Afghanistan  1800    3280000  Asia  Southern Asia  Low
1  Afghanistan  1801    3280000  Asia  Southern Asia  Low
2  Afghanistan  1802    3280000  Asia  Southern Asia  Low
3  Afghanistan  1803    3280000  Asia  Southern Asia  Low
4  Afghanistan  1804    3280000  Asia  Southern Asia  Low
...
39151  Zimbabwe  1969    5010000  Africa  Sub-Saharan Africa  Low
39197  Zimbabwe  2015   15800000  Africa  Sub-Saharan Africa  Low
39198  Zimbabwe  2016   16200000  Africa  Sub-Saharan Africa  Low
39199  Zimbabwe  2017   16500000  Africa  Sub-Saharan Africa  Low
39200  Zimbabwe  2018   16900000  Africa  Sub-Saharan Africa  Low

   life_expectancy  income  children_per_woman  child_mortality \
0                28.2     603                 7.00           469.0
1                28.2     603                 7.00           469.0
2                28.2     603                 7.00           469.0
3                28.2     603                 7.00           469.0
4                28.2     603                 7.00           469.0

```

|       |      |      |      |       |
|-------|------|------|------|-------|
| ...   | ...  | ...  | ...  | ...   |
| 39151 | 57.2 | 2160 | 7.42 | 115.0 |
| 39197 | 58.3 | 1890 | 3.84 | 59.9  |
| 39198 | 59.3 | 1860 | 3.76 | 56.4  |
| 39199 | 59.8 | 1910 | 3.68 | 56.8  |
| 39200 | 60.2 | 1950 | 3.61 | 55.5  |

|   | pop_density | co2_per_capita | years_in_school_men | years_in_school_women |
|---|-------------|----------------|---------------------|-----------------------|
| 0 | NaN         | NaN            | NaN                 | NaN                   |
| 1 | NaN         | NaN            | NaN                 | NaN                   |
| 2 | NaN         | NaN            | NaN                 | NaN                   |
| 3 | NaN         | NaN            | NaN                 | NaN                   |
| 4 | NaN         | NaN            | NaN                 | NaN                   |

|       |      |      |      |      |
|-------|------|------|------|------|
| ...   | ...  | ...  | ...  | ...  |
| 39151 | 12.9 | 1.35 | NaN  | NaN  |
| 39197 | 40.8 | NaN  | 11.1 | 10.2 |
| 39198 | 41.7 | NaN  | NaN  | NaN  |
| 39199 | 42.7 | NaN  | NaN  | NaN  |
| 39200 | 43.7 | NaN  | NaN  | NaN  |

[31484 rows x 14 columns]

Let's have a look at the list of unique countries and years whose rows include null values for the pop\_densityColumn:

```
[10]: print(set(data[data["pop_density"].isna()]["country"]))
print(set(data[data["pop_density"].isna()]["year"]))
```

```
{'Gambia', 'Montenegro', 'Swaziland', 'Iraq', 'Peru', 'Tunisia', 'Congo, Dem.
Rep.', 'Zimbabwe', 'Uganda', 'Chile', 'Czech Republic', 'Oman', 'Malaysia',
'Lithuania', 'Australia', 'Mauritania', 'Estonia', 'Poland', 'Netherlands',
'Bolivia', 'Sri Lanka', 'Benin', 'Italy', 'Niger', 'Slovenia', 'Cambodia',
'Afghanistan', 'Solomon Islands', 'Spain', 'Malawi', 'Gabon', 'Ghana', 'Canada',
'Denmark', 'Latvia', 'Norway', 'France', 'Equatorial Guinea', 'Jamaica',
'Guinea-Bissau', 'Senegal', 'Moldova', 'Burundi', 'Bahamas', 'Burkina Faso',
'Argentina', 'Liberia', 'South Sudan', 'Venezuela', 'Kiribati', 'Mexico',
'Switzerland', 'El Salvador', 'New Zealand', 'Sweden', 'Kazakhstan', 'South
Africa', 'Samoa', 'Croatia', 'Djibouti', 'Trinidad and Tobago', 'Brazil',
'Thailand', 'Tonga', 'Kenya', 'Slovak Republic', 'Ireland', 'Central African
Republic', 'Bhutan', 'Lebanon', 'Panama', 'Cuba', 'Serbia', 'Uzbekistan', 'Papua
New Guinea', 'Greece', 'Bulgaria', 'Bangladesh', 'Saudi Arabia', 'Lao',
'Mongolia', 'Ecuador', 'Cyprus', 'United Arab Emirates', 'Somalia', 'Vietnam',
'Seychelles', 'Paraguay', 'Palestine', 'Romania', 'Libya', 'Botswana',
'Algeria', 'Germany', 'Zambia', 'Tajikistan', 'Azerbaijan', 'Georgia',
'Ethiopia', 'Macedonia, FYR', 'United Kingdom', 'Uruguay', 'Suriname', 'Costa
Rica', 'Egypt', 'Guinea', 'Pakistan', 'Hungary', 'Rwanda', 'Mozambique',
'Mauritius', 'Kuwait', 'Togo', 'Kyrgyz Republic', 'Syria', 'Iceland', 'Chad',
```

'Barbados', 'Qatar', 'Albania', 'Nepal', 'Namibia', 'Dominican Republic', 'Guatemala', 'Cote d'Ivoire', 'India', 'Bahrain', 'Belize', 'Nicaragua', 'Comoros', 'Finland', 'Malta', 'United States', 'Grenada', 'Israel', 'Nigeria', 'China', 'Armenia', 'Luxembourg', 'Japan', 'Tanzania', 'Haiti', 'Morocco', 'Bosnia and Herzegovina', 'Congo, Rep.', 'Mali', 'Sierra Leone', 'South Korea', 'Singapore', 'Fiji', 'Myanmar', 'Turkey', 'Jordan', 'North Korea', 'Sudan', 'Honduras', 'Belarus', 'Austria', 'Yemen', 'Maldives', 'Eritrea', 'Angola', 'Colombia', 'Iran', 'Belgium', 'Indonesia', 'Portugal', 'Russia', 'Timor-Leste', 'Turkmenistan', 'Cameroon', 'Madagascar', 'Guyana', 'Lesotho', 'Vanuatu', 'Antigua and Barbuda', 'Philippines', 'Ukraine'}

{1800, 1801, 1802, 1803, 1804, 1805, 1806, 1807, 1808, 1809, 1810, 1811, 1812, 1813, 1814, 1815, 1816, 1817, 1818, 1819, 1820, 1821, 1822, 1823, 1824, 1825, 1826, 1827, 1828, 1829, 1830, 1831, 1832, 1833, 1834, 1835, 1836, 1837, 1838, 1839, 1840, 1841, 1842, 1843, 1844, 1845, 1846, 1847, 1848, 1849, 1850, 1851, 1852, 1853, 1854, 1855, 1856, 1857, 1858, 1859, 1860, 1861, 1862, 1863, 1864, 1865, 1866, 1867, 1868, 1869, 1870, 1871, 1872, 1873, 1874, 1875, 1876, 1877, 1878, 1879, 1880, 1881, 1882, 1883, 1884, 1885, 1886, 1887, 1888, 1889, 1890, 1891, 1892, 1893, 1894, 1895, 1896, 1897, 1898, 1899, 1900, 1901, 1902, 1903, 1904, 1905, 1906, 1907, 1908, 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916, 1917, 1918, 1919, 1920, 1921, 1922, 1923, 1924, 1925, 1926, 1927, 1928, 1929, 1930, 1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942, 1943, 1944, 1945, 1946, 1947, 1948, 1949}

It makes sense to see some of these countries and years as some of them declared their independence recently, others underwent foreign military occupations or wars.

Next, Let's follow a similar approach for the 'co2\_per\_capita' column:

```
[11]: data[data["co2_per_capita"].isna()].sample(n=10)
```

```
[11]:
```

|       | country  | year | population | region   | sub_region \                    |
|-------|----------|------|------------|----------|---------------------------------|
| 31094 | Slovenia | 2015 | 2070000    | Europe   | Southern Europe                 |
| 38406 | Vietnam  | 1881 | 11600000   | Asia     | South-eastern Asia              |
| 19826 | Liberia  | 1916 | 702000     | Africa   | Sub-Saharan Africa              |
| 19511 | Lesotho  | 1820 | 282000     | Africa   | Sub-Saharan Africa              |
| 32916 | Suriname | 1866 | 114000     | Americas | Latin America and the Caribbean |
| 17626 | Jordan   | 1906 | 342000     | Asia     | Western Asia                    |
| 13387 | Greece   | 1828 | 2490000    | Europe   | Southern Europe                 |
| 38596 | Yemen    | 1852 | 2750000    | Asia     | Western Asia                    |
| 16250 | Iraq     | 1844 | 1310000    | Asia     | Western Asia                    |
| 15019 | Honduras | 1927 | 942000     | Americas | Latin America and the Caribbean |

|       | income_group | life_expectancy | income | children_per_woman \ |
|-------|--------------|-----------------|--------|----------------------|
| 31094 | High         | 80.7            | 29100  | 1.61                 |
| 38406 | Lower middle | 31.7            | 917    | 4.70                 |
| 19826 | Low          | 34.1            | 1030   | 6.19                 |
| 19511 | Lower middle | 32.8            | 398    | 5.84                 |
| 32916 | Upper middle | 32.9            | 2240   | 6.58                 |



|       |              |      |      |      |
|-------|--------------|------|------|------|
| 17626 | Upper middle | 32.1 | 1590 | 6.97 |
| 13387 | High         | 36.6 | 1520 | 6.03 |
| 38596 | Low          | 23.4 | 1020 | 6.88 |
| 16250 | Upper middle | 31.2 | 1070 | 7.13 |
| 15019 | Lower middle | 35.7 | 2650 | 6.35 |

|       | child_mortality | pop_density | co2_per_capita | years_in_school_men | \ |
|-------|-----------------|-------------|----------------|---------------------|---|
| 31094 | 2.4             | 103.0       | NaN            | 13.4                |   |
| 38406 | 417.0           | NaN         | NaN            | NaN                 |   |
| 19826 | 416.0           | NaN         | NaN            | NaN                 |   |
| 19511 | 407.0           | NaN         | NaN            | NaN                 |   |
| 32916 | 406.0           | NaN         | NaN            | NaN                 |   |
| 17626 | 417.0           | NaN         | NaN            | NaN                 |   |
| 13387 | 361.0           | NaN         | NaN            | NaN                 |   |
| 38596 | 540.0           | NaN         | NaN            | NaN                 |   |
| 16250 | 428.0           | NaN         | NaN            | NaN                 |   |
| 15019 | 371.0           | NaN         | NaN            | NaN                 |   |

|       | years_in_school_women |
|-------|-----------------------|
| 31094 | 14.2                  |
| 38406 | NaN                   |
| 19826 | NaN                   |
| 19511 | NaN                   |
| 32916 | NaN                   |
| 17626 | NaN                   |
| 13387 | NaN                   |
| 38596 | NaN                   |
| 16250 | NaN                   |
| 15019 | NaN                   |

```
[12]: print(set(data[data["co2_per_capita"].isna()]["country"]))
      print(set(data[data["co2_per_capita"].isna()]["year"]))
```

```
{'Gambia', 'Montenegro', 'Swaziland', 'Iraq', 'Peru', 'Tunisia', 'Congo, Dem. Rep.', 'Zimbabwe', 'Uganda', 'Chile', 'Czech Republic', 'Oman', 'Malaysia', 'Lithuania', 'Australia', 'Mauritania', 'Estonia', 'Poland', 'Netherlands', 'Bolivia', 'Sri Lanka', 'Benin', 'Italy', 'Niger', 'Slovenia', 'Cambodia', 'Afghanistan', 'Solomon Islands', 'Spain', 'Malawi', 'Gabon', 'Ghana', 'Canada', 'Denmark', 'Latvia', 'Norway', 'France', 'Equatorial Guinea', 'Jamaica', 'Guinea-Bissau', 'Senegal', 'Moldova', 'Burundi', 'Bahamas', 'Burkina Faso', 'Argentina', 'Liberia', 'South Sudan', 'Venezuela', 'Kiribati', 'Mexico', 'Switzerland', 'El Salvador', 'New Zealand', 'Sweden', 'Kazakhstan', 'South Africa', 'Samoa', 'Croatia', 'Djibouti', 'Trinidad and Tobago', 'Brazil', 'Thailand', 'Tonga', 'Kenya', 'Slovak Republic', 'Ireland', 'Central African Republic', 'Bhutan', 'Lebanon', 'Panama', 'Cuba', 'Serbia', 'Uzbekistan', 'PapuaNew Guinea', 'Greece', 'Bulgaria', 'Bangladesh', 'Saudi Arabia', 'Lao',
```

'Mongolia', 'Ecuador', 'Cyprus', 'United Arab Emirates', 'Somalia', 'Vietnam', 'Seychelles', 'Paraguay', 'Palestine', 'Romania', 'Libya', 'Botswana', 'Algeria', 'Germany', 'Zambia', 'Tajikistan', 'Azerbaijan', 'Georgia', 'Ethiopia', 'Macedonia, FYR', 'United Kingdom', 'Uruguay', 'Suriname', 'Costa Rica', 'Egypt', 'Guinea', 'Pakistan', 'Hungary', 'Rwanda', 'Mozambique', 'Mauritius', 'Kuwait', 'Togo', 'Kyrgyz Republic', 'Syria', 'Iceland', 'Chad', 'Barbados', 'Qatar', 'Albania', 'Nepal', 'Namibia', 'Dominican Republic', 'Guatemala', 'Cote d'Ivoire', 'India', 'Bahrain', 'Belize', 'Nicaragua', 'Comoros', 'Finland', 'Malta', 'United States', 'Grenada', 'Israel', 'Nigeria', 'China', 'Armenia', 'Luxembourg', 'Japan', 'Tanzania', 'Haiti', 'Morocco', 'Bosnia and Herzegovina', 'Congo, Rep.', 'Mali', 'Sierra Leone', 'South Korea', 'Singapore', 'Fiji', 'Myanmar', 'Turkey', 'Jordan', 'North Korea', 'Sudan', 'Honduras', 'Belarus', 'Austria', 'Yemen', 'Maldives', 'Eritrea', 'Angola', 'Colombia', 'Iran', 'Belgium', 'Indonesia', 'Portugal', 'Russia', 'Timor-Leste', 'Turkmenistan', 'Cameroon', 'Madagascar', 'Guyana', 'Lesotho', 'Vanuatu', 'Antigua and Barbuda', 'Philippines', 'Ukraine'}

{1800, 1801, 1802, 1803, 1804, 1805, 1806, 1807, 1808, 1809, 1810, 1811, 1812, 1813, 1814, 1815, 1816, 1817, 1818, 1819, 1820, 1821, 1822, 1823, 1824, 1825, 1826, 1827, 1828, 1829, 1830, 1831, 1832, 1833, 1834, 1835, 1836, 1837, 1838, 1839, 1840, 1841, 1842, 1843, 1844, 1845, 1846, 1847, 1848, 1849, 1850, 1851, 1852, 1853, 1854, 1855, 1856, 1857, 1858, 1859, 1860, 1861, 1862, 1863, 1864, 1865, 1866, 1867, 1868, 1869, 1870, 1871, 1872, 1873, 1874, 1875, 1876, 1877, 1878, 1879, 1880, 1881, 1882, 1883, 1884, 1885, 1886, 1887, 1888, 1889, 1890, 1891, 1892, 1893, 1894, 1895, 1896, 1897, 1898, 1899, 1900, 1901, 1902, 1903, 1904, 1905, 1906, 1907, 1908, 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916, 1917, 1918, 1919, 1920, 1921, 1922, 1923, 1924, 1925, 1926, 1927, 1928, 1929, 1930, 1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942, 1943, 1944, 1945, 1946, 1947, 1948, 1949, 1950, 1951, 1952, 1953, 1954, 1955, 1956, 1957, 1958, 1959, 1960, 1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, 1970, 1971, 1972, 1973, 1974, 1975, 1976, 1977, 1978, 1979, 1980, 1981, 1982, 1983, 1984, 1985, 1986, 1987, 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2015, 2016, 2017, 2018}

We can also see countries that have undergone difficult periods. Given the complexity of these events, it will be difficult to make predictions or find the right value to replace for all countries, i.e. the mean value might be good for some but it might lead to uncertainty in others. Furthermore, we can see that there is a period (2013-2015) in which there are no null values. Therefore, it seems like the best approach is to **remove the null values**.

```
[13]: data = data.dropna().reset_index(drop=True)
      data.isnull().sum()
```

```
[13]: country      0
      year         0
      population   0
      region       0
```

```

sub_region      0
income_group    0
life_expectancy 0
income          0
children_per_woman 0
child_mortality 0
pop_density     0
co2_per_capita  0
years_in_school_men 0
years_in_school_women 0
dtype: int64

```

## 4.2 Select a subset

Given the range of the data (1800-2018), drawing conclusions based on the entire data set is difficult and counter productive. We can expect similar relationship for adjacent years, such as 1995-1998, 1998-2000, 2005-2007 and so on, but most likely those relationship have changed between 1800 and 2018. Therefore, working with the entire data set might lead to not seeing the full picture.

This being said, we can select a subset of the data and in the data cleaning section we identified a subset that included no null values and it isn't affected by the loss of information.

```

[14]: # select the subset belonging to the year range 2013-2015
data_period = data[
    (data["year"] == 2013) | (data["year"] == 2014) | (data["year"] == 2015)
]
data_period.head()

```

```

[14]:
   country  year  population  region  sub_region  income_group \
43  Afghanistan  2013   31700000  Asia  Southern Asia  Low
44  Afghanistan  2014   32800000  Asia  Southern Asia  Low
88    Albania  2013    2920000  Europe  Southern Europe  Upper middle
89    Albania  2014    2920000  Europe  Southern Europe  Upper middle
133   Algeria  2013   38300000  Africa  Northern Africa  Upper middle

   life_expectancy  income  children_per_woman  child_mortality \
43      57.7      1810      5.17      79.3
44      57.8      1780      4.98      76.1
88      77.2     10500      1.70      14.9
89      77.4     10700      1.71      14.4
133     77.0     13300      2.92      25.8

   pop_density  co2_per_capita  years_in_school_men  years_in_school_women
43      48.6      0.316      3.94      0.92
44      50.2      0.299      4.04      0.95
88     107.0      1.730     11.70     11.90
89     107.0      1.960     11.80     12.10

```

133            16.1            3.510            8.24            7.42

[15]: data\_period.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 358 entries, 43 to 7716

Data columns (total 14 columns):

| #  | Column                | Non-Null Count | Dtype   |
|----|-----------------------|----------------|---------|
| 0  | country               | 358 non-null   | object  |
| 1  | year                  | 358 non-null   | int64   |
| 2  | population            | 358 non-null   | int64   |
| 3  | region                | 358 non-null   | object  |
| 4  | sub_region            | 358 non-null   | object  |
| 5  | income_group          | 358 non-null   | object  |
| 6  | life_expectancy       | 358 non-null   | float64 |
| 7  | income                | 358 non-null   | int64   |
| 8  | children_per_woman    | 358 non-null   | float64 |
| 9  | child_mortality       | 358 non-null   | float64 |
| 10 | pop_density           | 358 non-null   | float64 |
| 11 | co2_per_capita        | 358 non-null   | float64 |
| 12 | years_in_school_men   | 358 non-null   | float64 |
| 13 | years_in_school_women | 358 non-null   | float64 |

dtypes: float64(7), int64(3), object(4)

memory usage: 42.0+ KB

[16]: data\_period.groupby("region").describe()

```
[16]:
```

|          | year  | count  | mean     | std    | min    | 25%    | 50%    | 75%    | max |
|----------|-------|--------|----------|--------|--------|--------|--------|--------|-----|
| region   |       |        |          |        |        |        |        |        |     |
| Africa   | 104.0 | 2013.5 | 0.502421 | 2013.0 | 2013.0 | 2013.5 | 2014.0 | 2014.0 |     |
| Americas | 62.0  | 2013.5 | 0.504082 | 2013.0 | 2013.0 | 2013.5 | 2014.0 | 2014.0 |     |
| Asia     | 94.0  | 2013.5 | 0.502681 | 2013.0 | 2013.0 | 2013.5 | 2014.0 | 2014.0 |     |
| Europe   | 78.0  | 2013.5 | 0.503236 | 2013.0 | 2013.0 | 2013.5 | 2014.0 | 2014.0 |     |
| Oceania  | 20.0  | 2013.5 | 0.512989 | 2013.0 | 2013.0 | 2013.5 | 2014.0 | 2014.0 |     |

|          | population | count        | mean | ... | years_in_school_men | 75%  | max |
|----------|------------|--------------|------|-----|---------------------|------|-----|
| region   |            |              |      | ... |                     |      |     |
| Africa   | 104.0      | 2.205035e+07 | ...  |     | 8.9275              | 11.3 |     |
| Americas | 62.0       | 3.124632e+07 | ...  |     | 11.8750             | 15.3 |     |
| Asia     | 94.0       | 9.190187e+07 | ...  |     | 12.4000             | 15.0 |     |
| Europe   | 78.0       | 1.896846e+07 | ...  |     | 13.9000             | 14.8 |     |
| Oceania  | 20.0       | 3.546855e+07 | ...  |     | 13.9250             | 14.5 |     |

|  | years_in_school_women |
|--|-----------------------|
|  |                       |

| region   | count | mean      | std      | min   | 25%     | 50%   |
|----------|-------|-----------|----------|-------|---------|-------|
| Africa   | 104.0 | 5.646827  | 2.886351 | 1.12  | 3.5250  | 5.38  |
| Americas | 62.0  | 11.053226 | 2.121167 | 6.35  | 9.7200  | 10.85 |
| Asia     | 94.0  | 9.859362  | 3.658881 | 0.92  | 7.5700  | 10.60 |
| Europe   | 78.0  | 13.733333 | 1.039938 | 11.00 | 13.3250 | 14.00 |
| Oceania  | 20.0  | 11.268000 | 3.093146 | 5.75  | 8.7375  | 11.90 |

| region   | 75%    | max  |
|----------|--------|------|
| Africa   | 7.795  | 11.4 |
| Americas | 12.700 | 15.5 |
| Asia     | 12.800 | 15.6 |
| Europe   | 14.400 | 15.3 |
| Oceania  | 14.325 | 14.9 |

[5 rows x 80 columns]

## 5 Outliers

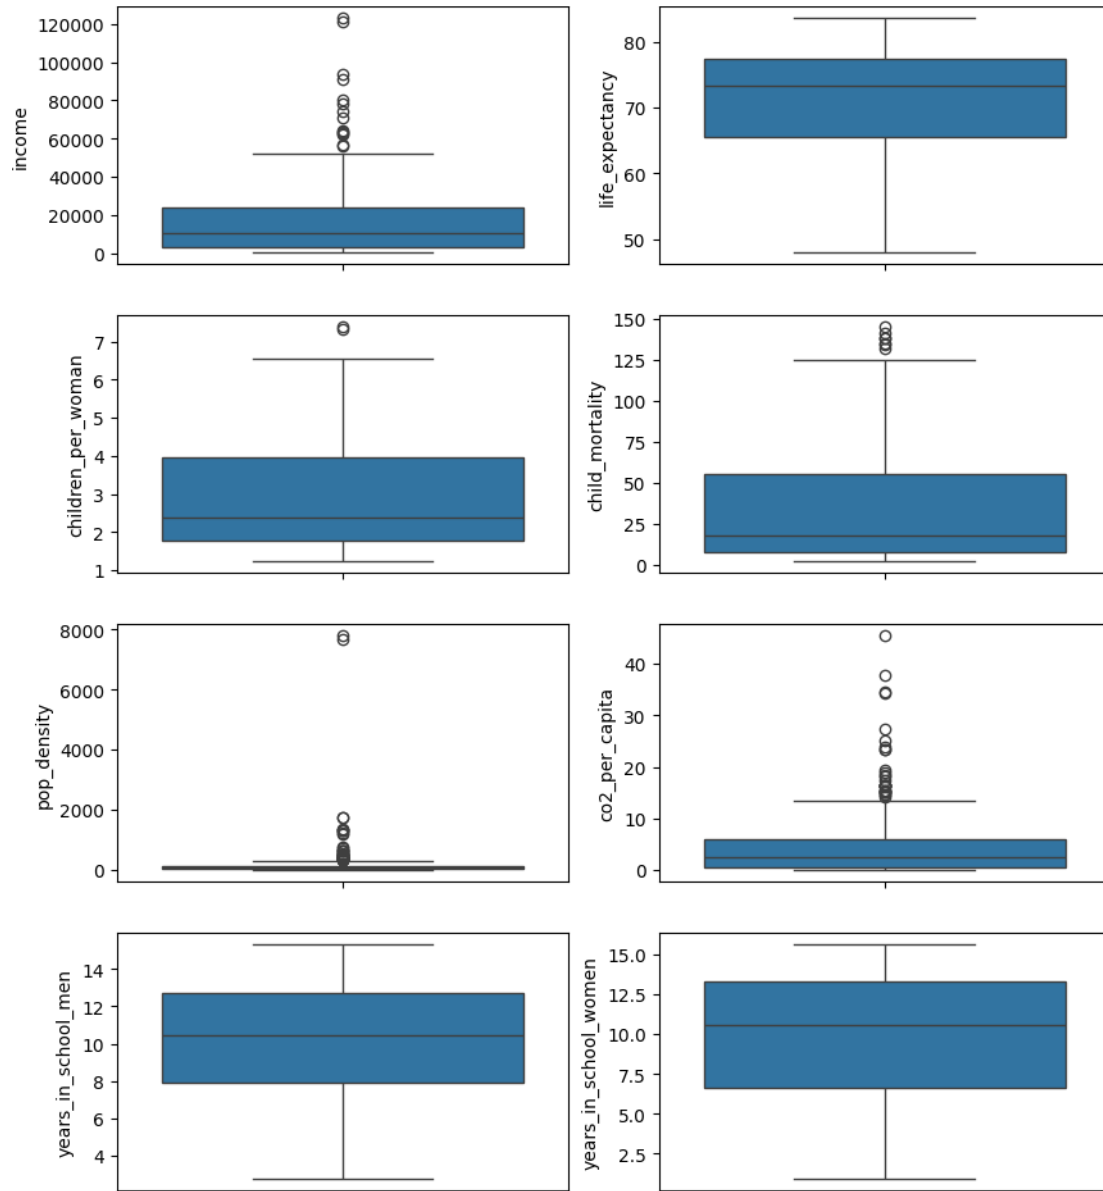
The presence of outliers can easily be identified by using box plots. However, one has to be careful as a boxplot of the entire data frame will show many outliers. Let's create the box plots

```
[17]: # Create a list using the features that contain numerical data
column_list_plot = [
    "income",
    "life_expectancy",
    "children_per_woman",
    "child_mortality",
    "pop_density",
    "co2_per_capita",
    "years_in_school_men",
    "years_in_school_women",
]

# Create a 4x2 figure with 8 subplots, where 8 of them will be used
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(10, 12))

# Flatten the axes array to iterate over it
axes = axes.flatten()

# Use a for loop to create the subplots
for index, col_name in enumerate(column_list_plot):
    sns.boxplot(ax=axes[index], y=col_name, data=data_period)
```



As expected, the box plots for four of the features (income, child\_mortality, pop\_density and co2\_emissions) show many outliers, but this isn't the full picture. An accurate analysis of the outliers requires to plot the data split into continents.

```
[18]: # Create a list of colors for different regions
region_colors = {
    "Africa": "green",
    "Americas": "yellow",
    "Asia": "blue",
    "Europe": "purple",
    "Oceania": "red",
```

```

}

# Create a 4x2 figure with 8 subplots, Where 8 of them will be used
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(12, 15))

# Flatten the axes array to iterate over it
axes = axes.flatten()

# Use a for loop to create the subplots
for index, col_name in enumerate(column_list_plot):
    sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
               palette=region_colors)

# Add a title to the entire figure
fig.suptitle("Box Plots of Numerical Features by Region", fontsize=20)

```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
```

```
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

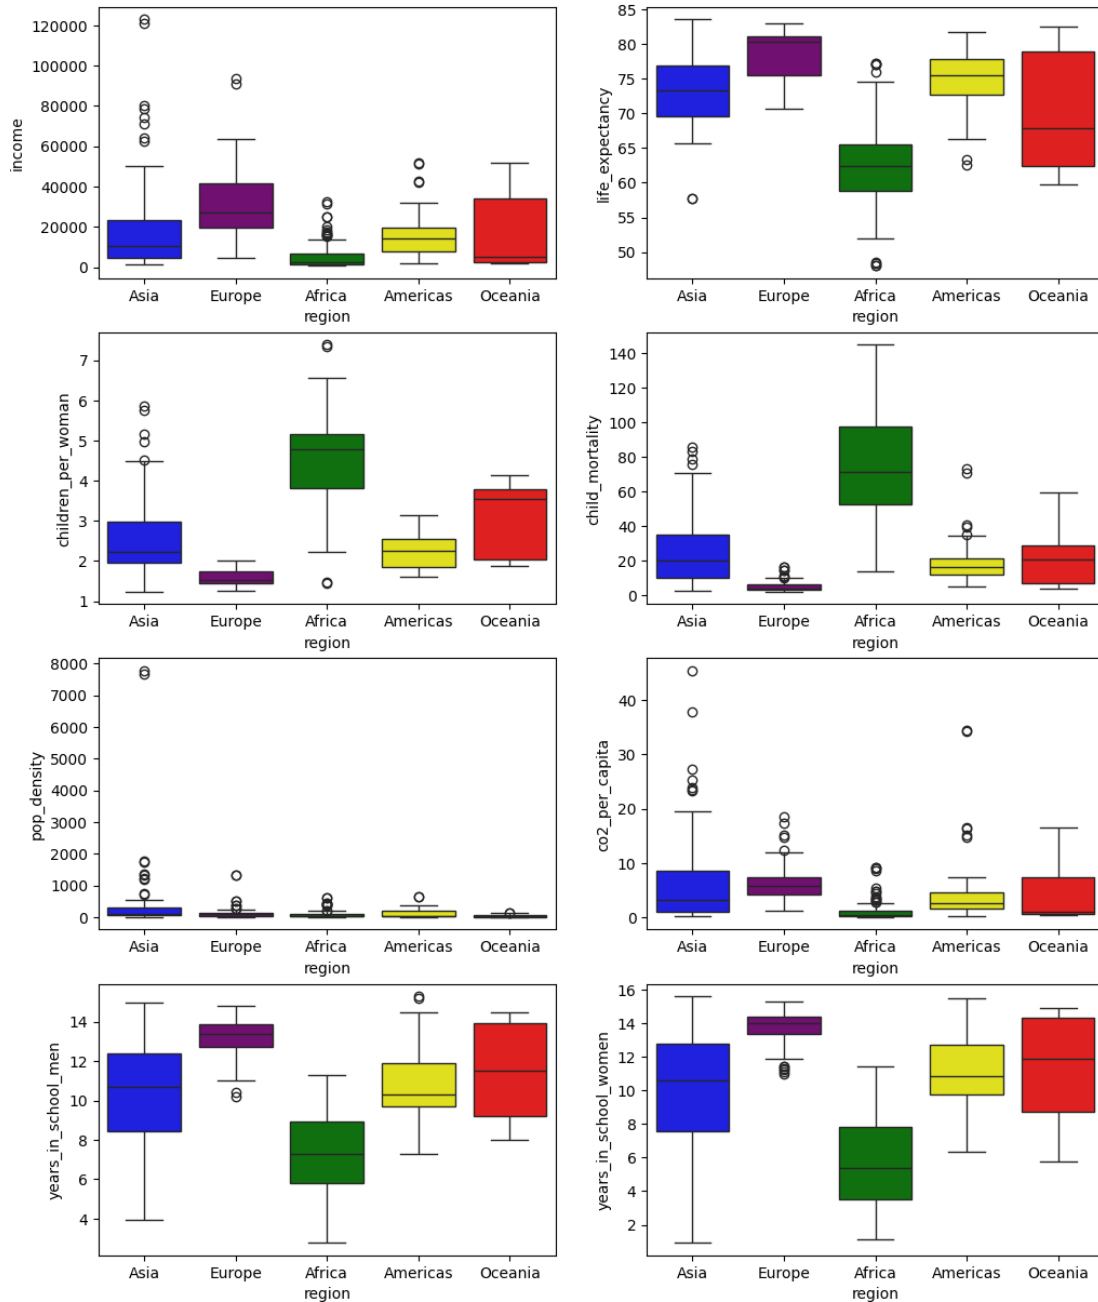
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

[18]: Text(0.5, 0.98, 'Box Plots of Numerical Features by Region')



## Box Plots of Numerical Features by Region



Actually, these box plots show more outliers than the previous graphs. As studied in this course, the presence of outliers can be due to errors in the measurement, but it can also be right. Let's have a look at data belonging to Oceania to understand the reason behind the outliers:

```
[19]: data_period[data_period["region"] == "Oceania"]
```

```
[19]:
```

|      | country          | year | population | region  | sub_region \              |
|------|------------------|------|------------|---------|---------------------------|
| 358  | Australia        | 2013 | 23200000   | Oceania | Australia and New Zealand |
| 359  | Australia        | 2014 | 23500000   | Oceania | Australia and New Zealand |
| 2336 | Fiji             | 2013 | 880000     | Oceania | Melanesia                 |
| 2337 | Fiji             | 2014 | 886000     | Oceania | Melanesia                 |
| 3686 | Kiribati         | 2013 | 109000     | Oceania | Micronesia                |
| 3687 | Kiribati         | 2014 | 110000     | Oceania | Micronesia                |
| 4928 | New Zealand      | 2013 | 4520000    | Oceania | Australia and New Zealand |
| 4929 | New Zealand      | 2014 | 4570000    | Oceania | Australia and New Zealand |
| 5358 | Papua New Guinea | 2013 | 7590000    | Oceania | Melanesia                 |
| 5359 | Papua New Guinea | 2014 | 7760000    | Oceania | Melanesia                 |
| 5808 | Samoa            | 2013 | 191000     | Oceania | Polynesia                 |
| 5809 | Samoa            | 2014 | 192000     | Oceania | Polynesia                 |
| 6169 | Solomon Islands  | 2013 | 564000     | Oceania | Melanesia                 |
| 6170 | Solomon Islands  | 2014 | 576000     | Oceania | Melanesia                 |
| 6905 | Tonga            | 2013 | 105000     | Oceania | Polynesia                 |
| 6906 | Tonga            | 2014 | 106000     | Oceania | Polynesia                 |
| 7354 | United States    | 2013 | 316000000  | Oceania | Micronesia                |
| 7356 | United States    | 2014 | 318000000  | Oceania | Micronesia                |
| 7490 | Vanuatu          | 2013 | 253000     | Oceania | Melanesia                 |
| 7491 | Vanuatu          | 2014 | 259000     | Oceania | Melanesia                 |

|      | income_group | life_expectancy | income | children_per_woman \ |
|------|--------------|-----------------|--------|----------------------|
| 358  | High         | 82.5            | 42900  | 1.89                 |
| 359  | High         | 82.6            | 43400  | 1.87                 |
| 2336 | Upper middle | 65.5            | 7980   | 2.59                 |
| 2337 | Upper middle | 65.5            | 8350   | 2.57                 |
| 3686 | Lower middle | 61.2            | 1830   | 3.77                 |
| 3687 | Lower middle | 61.4            | 1840   | 3.73                 |
| 4928 | High         | 81.5            | 33800  | 2.05                 |
| 4929 | High         | 81.5            | 34500  | 2.03                 |
| 5358 | Lower middle | 59.8            | 2470   | 3.81                 |
| 5359 | Lower middle | 60.1            | 2620   | 3.76                 |
| 5808 | Upper middle | 71.6            | 5490   | 4.15                 |
| 5809 | Upper middle | 71.6            | 5510   | 4.09                 |
| 6169 | Lower middle | 62.4            | 2030   | 4.03                 |
| 6170 | Lower middle | 62.4            | 2020   | 3.97                 |
| 6905 | Upper middle | 70.1            | 4950   | 3.77                 |
| 6906 | Upper middle | 70.2            | 5030   | 3.72                 |
| 7354 | High         | 78.9            | 51000  | 1.96                 |
| 7356 | High         | 78.9            | 51800  | 1.95                 |
| 7490 | Lower middle | 63.5            | 2890   | 3.38                 |
| 7491 | Lower middle | 63.5            | 2890   | 3.35                 |

| child_mortality | pop_density | co2_per_capita | years_in_school_men \ |
|-----------------|-------------|----------------|-----------------------|
|-----------------|-------------|----------------|-----------------------|

|      |      |        |        |       |
|------|------|--------|--------|-------|
| 358  | 4.2  | 3.01   | 16.100 | 13.90 |
| 359  | 4.0  | 3.06   | 15.400 | 14.00 |
| 2336 | 23.4 | 48.20  | 1.310  | 11.40 |
| 2337 | 23.0 | 48.50  | 1.320  | 11.50 |
| 3686 | 58.8 | 134.00 | 0.574  | 9.36  |
| 3687 | 57.4 | 136.00 | 0.564  | 9.48  |
| 4928 | 5.9  | 17.20  | 7.410  | 14.20 |
| 4929 | 5.7  | 17.30  | 7.590  | 14.30 |
| 5358 | 59.5 | 16.80  | 0.815  | 8.01  |
| 5359 | 57.9 | 17.10  | 0.815  | 8.15  |
| 5808 | 18.4 | 67.40  | 1.040  | 12.10 |
| 5809 | 18.1 | 67.90  | 1.030  | 12.20 |
| 6169 | 27.2 | 20.10  | 0.358  | 8.63  |
| 6170 | 26.8 | 20.60  | 0.350  | 8.78  |
| 6905 | 17.4 | 146.00 | 1.080  | 11.50 |
| 6906 | 17.1 | 147.00 | 1.140  | 11.60 |
| 7354 | 6.9  | 34.50  | 16.400 | 14.40 |
| 7356 | 6.8  | 34.70  | 16.500 | 14.50 |
| 7490 | 29.1 | 20.80  | 0.420  | 9.09  |
| 7491 | 28.7 | 21.20  | 0.595  | 9.23  |

|      | years_in_school_women |
|------|-----------------------|
| 358  | 14.30                 |
| 359  | 14.40                 |
| 2336 | 11.80                 |
| 2337 | 11.90                 |
| 3686 | 9.70                  |
| 3687 | 9.85                  |
| 4928 | 14.80                 |
| 4929 | 14.90                 |
| 5358 | 5.75                  |
| 5359 | 5.89                  |
| 5808 | 12.90                 |
| 5809 | 13.00                 |
| 6169 | 7.50                  |
| 6170 | 7.68                  |
| 6905 | 11.90                 |
| 6906 | 12.00                 |
| 7354 | 14.80                 |
| 7356 | 14.90                 |
| 7490 | 8.61                  |
| 7491 | 8.78                  |

We can see that the outliers from the `income` column are Australia and New Zealand. These values represent actual data and removing them would cause a loss of data. The same can be said regarding the `child_mortality`, `pop_density` and `co2_per_capita` besides any other feature for any other continents. Therefore, ***The outliers would not be removed.***

## 5.1 Exploratory Data Analysis

##Univariate Analysis

```
[20]: # For numerical attributes, let's describe the dataset to get mean, median, etc.  
data.describe()
```

```
[20]:
```

|       | year        | population   | life_expectancy | income \      |
|-------|-------------|--------------|-----------------|---------------|
| count | 7717.000000 | 7.717000e+03 | 7717.000000     | 7717.000000   |
| mean  | 1992.435921 | 3.313877e+07 | 66.028768       | 12962.267850  |
| std   | 12.979325   | 1.188602e+08 | 9.654891        | 17553.735789  |
| min   | 1970.000000 | 5.120000e+04 | 12.600000       | 247.000000    |
| 25%   | 1981.000000 | 2.440000e+06 | 58.700000       | 2280.000000   |
| 50%   | 1993.000000 | 7.120000e+06 | 68.400000       | 6560.000000   |
| 75%   | 2004.000000 | 2.080000e+07 | 73.600000       | 16600.000000  |
| max   | 2014.000000 | 1.390000e+09 | 83.600000       | 178000.000000 |


  

|       | children_per_woman | child_mortality | pop_density | co2_per_capita \ |
|-------|--------------------|-----------------|-------------|------------------|
| count | 7717.000000        | 7717.000000     | 7717.000000 | 7717.000000      |
| mean  | 3.913842           | 74.726785       | 136.687679  | 4.665658         |
| std   | 1.990729           | 73.384581       | 417.938633  | 7.215037         |
| min   | 1.120000           | 2.300000        | 0.823000    | 0.004330         |
| 25%   | 2.060000           | 17.000000       | 18.000000   | 0.421000         |
| 50%   | 3.480000           | 46.700000       | 53.400000   | 1.870000         |
| 75%   | 5.710000           | 113.000000      | 122.000000  | 6.570000         |
| max   | 8.870000           | 399.000000      | 7780.000000 | 87.700000        |

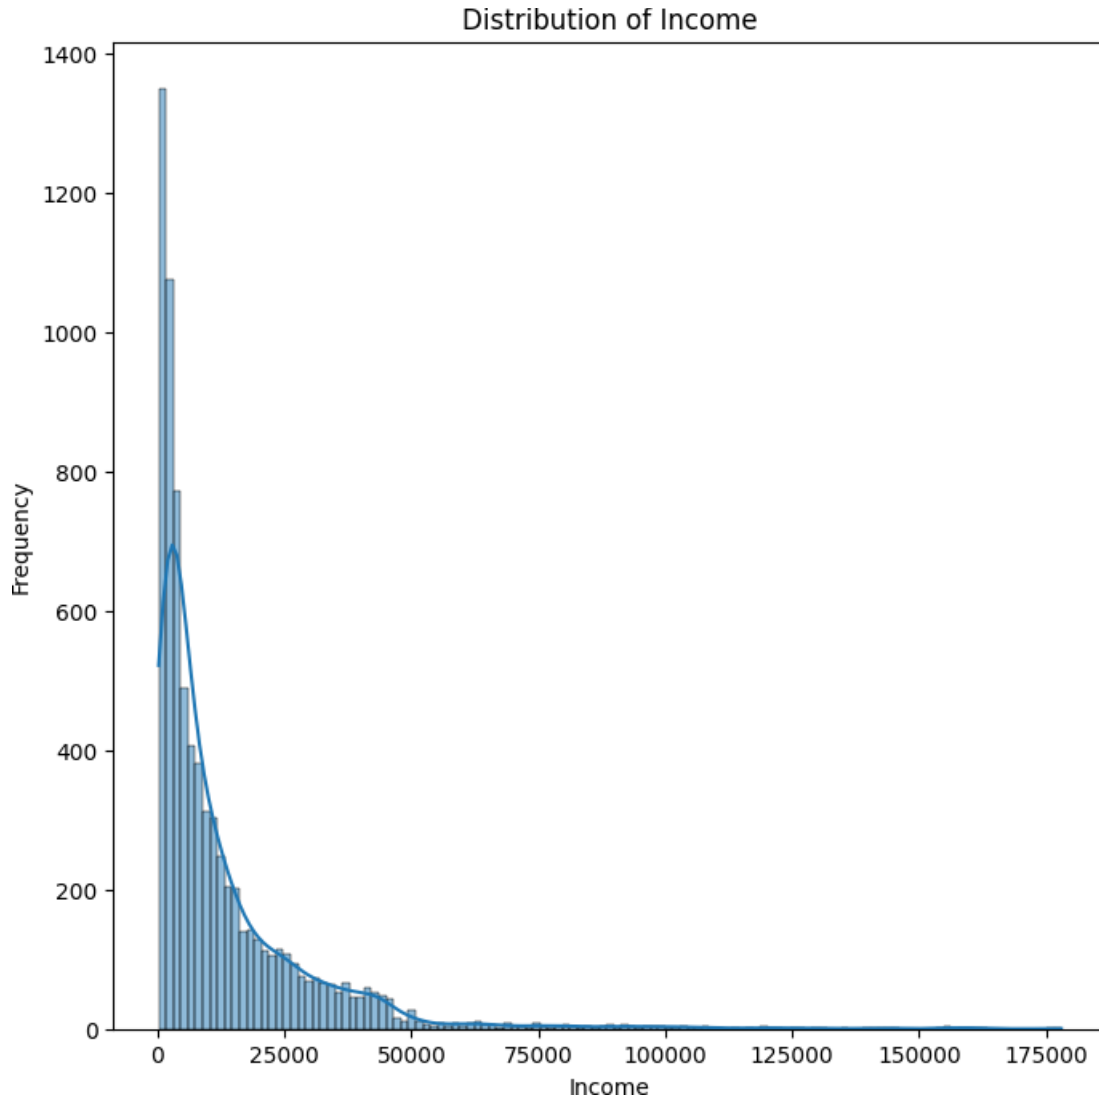
|       | years_in_school_men | years_in_school_women |
|-------|---------------------|-----------------------|
| count | 7717.000000         | 7717.000000           |
| mean  | 7.720621            | 6.981117              |
| std   | 3.190283            | 3.888600              |
| min   | 0.900000            | 0.210000              |
| 25%   | 5.180000            | 3.600000              |
| 50%   | 7.680000            | 7.030000              |
| 75%   | 10.200000           | 10.100000             |
| max   | 15.300000           | 15.600000             |

```
[21]: # Visualizing distributions of numerical attributes  
# Histogram for Life Expectancy  
  
plt.figure(figsize=(10, 6))  
sns.histplot(data["life_expectancy"], kde=True)  
plt.title("Distribution of Life Expectancy")  
plt.xlabel("Life Expectancy")  
plt.ylabel("Frequency")  
plt.show()
```



Here, we can see Life Expectancy is left-skewed, which means that most of the data points are concentrated towards the higher end of the scale, while a few extreme values (outliers) pull the distribution towards the lower end. In other words, the majority of countries or regions tend to have relatively high life expectancies, but there are a few places with significantly lower life expectancies that drag the overall distribution to the left. This skewness can occur due to various factors, such as differences in healthcare, socioeconomic conditions, and lifestyle choices etc.

```
[22]: # Histogram for Income
plt.figure(figsize=(8, 8))
sns.histplot(data["income"], kde=True) # Replace 'income' with the relevant_
column name
plt.title("Distribution of Income")
plt.xlabel("Income")
plt.ylabel("Frequency")
plt.show()
```



The right-skewed distribution of income indicates that most data points are concentrated towards the lower end of the income scale, while a few extreme values (outliers) pull the distribution towards the higher end. In other words, the majority of individuals tend to have relatively lower incomes, but there are a few high-income outliers that stretch the overall distribution to the right. Factors contributing to this skewness include income disparities, wealth concentration, and economic inequality.

## 5.2 Bivariate Analysis

1. Line Charts
2. Scatter Plots
3. Correlation plots


```
[23]: ## Line Charts or Visualizing the trend
life_expectancy_regional_data = data.groupby(['region', 'year'], as_index=False).agg({'life_expectancy': 'mean'})
print(life_expectancy_regional_data)
```

|     | region  | year | life_expectancy |
|-----|---------|------|-----------------|
| 0   | Africa  | 1970 | 51.029787       |
| 1   | Africa  | 1971 | 51.374468       |
| 2   | Africa  | 1972 | 51.364583       |
| 3   | Africa  | 1973 | 52.268750       |
| 4   | Africa  | 1974 | 52.650000       |
| ... | ...     | ...  | ...             |
| 220 | Oceania | 10   | 69.240000       |
| 221 | Oceania | 11   | 69.400000       |
| 222 | Oceania | 12   | 69.500000       |
| 223 | Oceania | 13   | 69.700000       |
| 224 | Oceania | 14   | 69.770000       |

[225 rows x 3 columns]

```
[24]: import plotly.express as px
import plotly.graph_objects as go

#Comparing Life expectancy among regions
fig = px.line(data_frame= life_expectancy_regional_data, x='year', y='life_expectancy', color='region', labels = {'year': 'Year', 'life_expectancy': 'Life Expectancy'}, title = '% Employed in the Life Expectancy Across Time')
fig.show()
```



Although Europe has the highest life\_expectancy index and Africa having the lowest index, Asia seems to have the highest rate of increase in Life expectancy!

[25]: **5.2.1 Let's now take a look at our Income Plots**

```
income_region_data = data.groupby(['region', 'year'], as_index=False).agg({'income': 'mean'})
print(income_region_data)
```

|     | region  | year | income       |
|-----|---------|------|--------------|
| 0   | Africa  | 1970 | 3866.170213  |
| 1   | Africa  | 1971 | 3826.659574  |
| 2   | Africa  | 1972 | 3676.395833  |
| 3   | Africa  | 1973 | 3712.062500  |
| 4   | Africa  | 1974 | 3791.562500  |
| ..  | ...     | ...  | ...          |
| 220 | Oceania | 2010 | 14938.000000 |
| 221 | Oceania | 2011 | 15142.000000 |
| 222 | Oceania | 2012 | 15381.000000 |
| 223 | Oceania | 2013 | 15534.000000 |
| 224 | Oceania | 2014 | 15796.000000 |

```
[26]: fig = px.line(data_frame=income_region_data, x='year', y='income',
↳color='region', labels = {'year': 'Year', 'income': 'Income'}, title = '%_
↳Employed in Income Across Time')
#add annotations
fig.show()
```





### 5.2.2 Now we are able to get ideas about our initial 2 questions: How have socio-economic indicators evolved globally over the years? Can we predict future trends in key indicators?

Which region has the highest Life Expectancy? Is it the same for Income? Europe has the highest Life Expectancy and Income value across time, and Africa at the lowest for both

### 5.3 What about Life expectancy and Income values for countries in each continent?

Let's start by grouping and aggregating our data! We will examine Life Expectancy first

```
life_exp_ctype_region_data = data.groupby(["country", "region", "year"], _
↳as_index= False).agg({"life_expectancy": "mean"})
print(life_exp_ctype_region_data)
```

|      | country     | region | year | life_expectancy |
|------|-------------|--------|------|-----------------|
| 0    | Afghanistan | Asia   | 1970 | 45.8            |
| 1    | Afghanistan | Asia   | 1971 | 45.9            |
| 2    | Afghanistan | Asia   | 1972 | 45.9            |
| 3    | Afghanistan | Asia   | 1973 | 46.0            |
| 4    | Afghanistan | Asia   | 1974 | 46.1            |
| ...  | ...         | ...    | ...  | ...             |
| 7712 | Zimbabwe    | Africa | 2010 | 49.6            |
| 7713 | Zimbabwe    | Africa | 2011 | 51.9            |
| 7714 | Zimbabwe    | Africa | 2012 | 54.1            |
| 7715 | Zimbabwe    | Africa | 2013 | 55.6            |
| 7716 | Zimbabwe    | Africa | 2014 | 57.0            |

[7717 rows x 4 columns]

#### Let's examine Americas' Life Expectancy

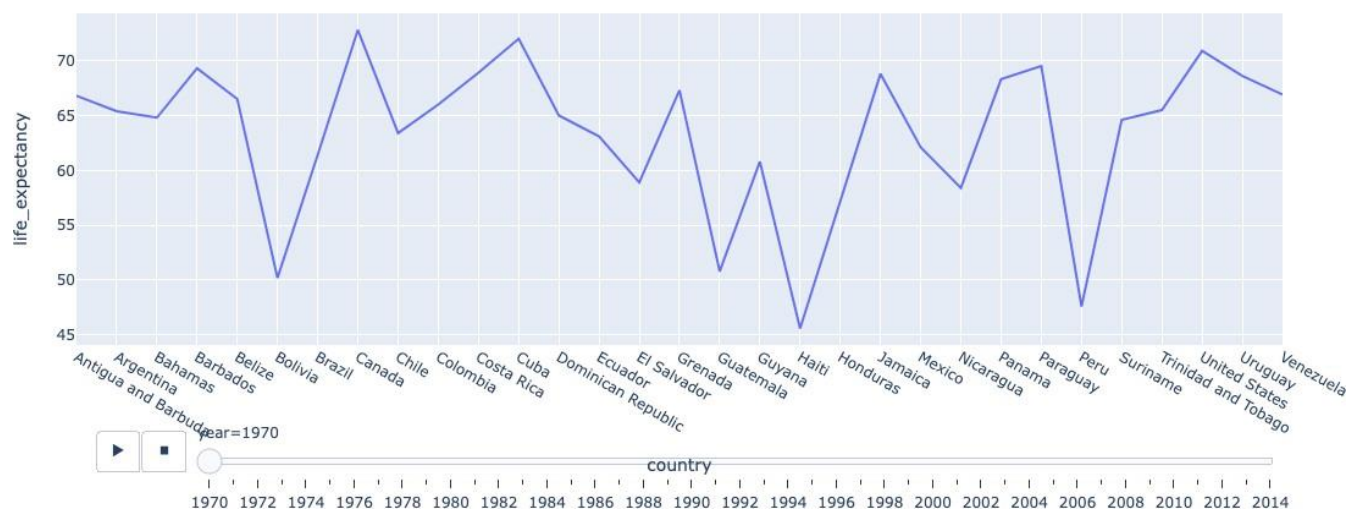
```
[28]: ame_lif_exp = life_exp_ctype_region_data[life_exp_ctype_region_data.
↳region=="Americas"]
```

```
print(ame_lif_exp)
```

|      | country             | region   | year | life_expectancy |
|------|---------------------|----------|------|-----------------|
| 180  | Antigua and Barbuda | Americas | 1970 | 66.8            |
| 181  | Antigua and Barbuda | Americas | 1971 | 67.2            |
| 182  | Antigua and Barbuda | Americas | 1972 | 67.6            |
| 183  | Antigua and Barbuda | Americas | 1973 | 68.0            |
| 184  | Antigua and Barbuda | Americas | 1974 | 68.3            |
| ...  | ...                 | ...      | ...  | ...             |
| 7532 | Venezuela           | Americas | 2010 | 75.4            |
| 7533 | Venezuela           | Americas | 2011 | 75.4            |
| 7534 | Venezuela           | Americas | 2012 | 75.3            |
| 7535 | Venezuela           | Americas | 2013 | 75.4            |
| 7536 | Venezuela           | Americas | 2014 | 75.5            |

[1395 rows x 4 columns]

Life Expectancy in Americas (PLZ USE AUTOSCALE FOR THE LINE TO SEE)



Life Expectancy Index for United States and Canada are the highest in Americas where Haiti has the least life expectancy

### 5.3.1 Examining Income data by Country

[30]:

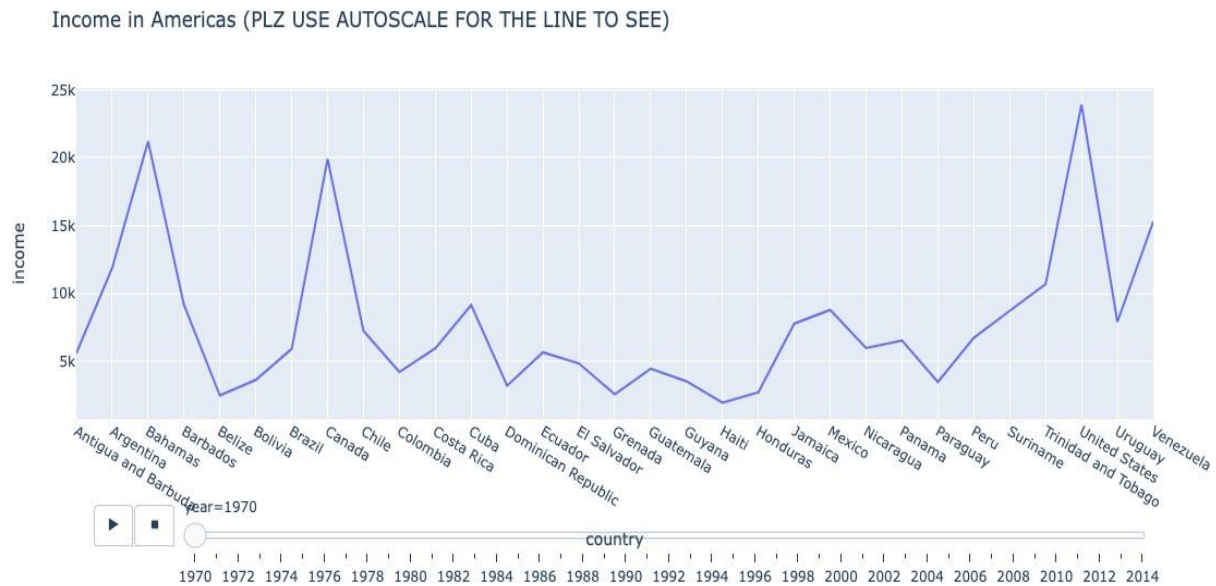
```
income_ctype_region_data = data.groupby(['country', 'region', 'year'], as_index=False).agg({'income': 'mean'})
income_ctype_region_data
```

|   | country     | region | year | income |
|---|-------------|--------|------|--------|
| 0 | Afghanistan | Asia   | 1970 | 1180.0 |
| 1 | Afghanistan | Asia   | 1971 | 1100.0 |

|      |             |        |      |        |
|------|-------------|--------|------|--------|
| 2    | Afghanistan | Asia   | 1972 | 1050.0 |
| 3    | Afghanistan | Asia   | 1973 | 1150.0 |
| 4    | Afghanistan | Asia   | 1974 | 1180.0 |
| ...  | ...         | ...    | ...  | ...    |
| 7712 | Zimbabwe    | Africa | 2010 | 1460.0 |
| 7713 | Zimbabwe    | Africa | 2011 | 1660.0 |
| 7714 | Zimbabwe    | Africa | 2012 | 1850.0 |
| 7715 | Zimbabwe    | Africa | 2013 | 1900.0 |
| 7716 | Zimbabwe    | Africa | 2014 | 1910.0 |

[7717 rows x 4 columns]

## Taking a look at America Again

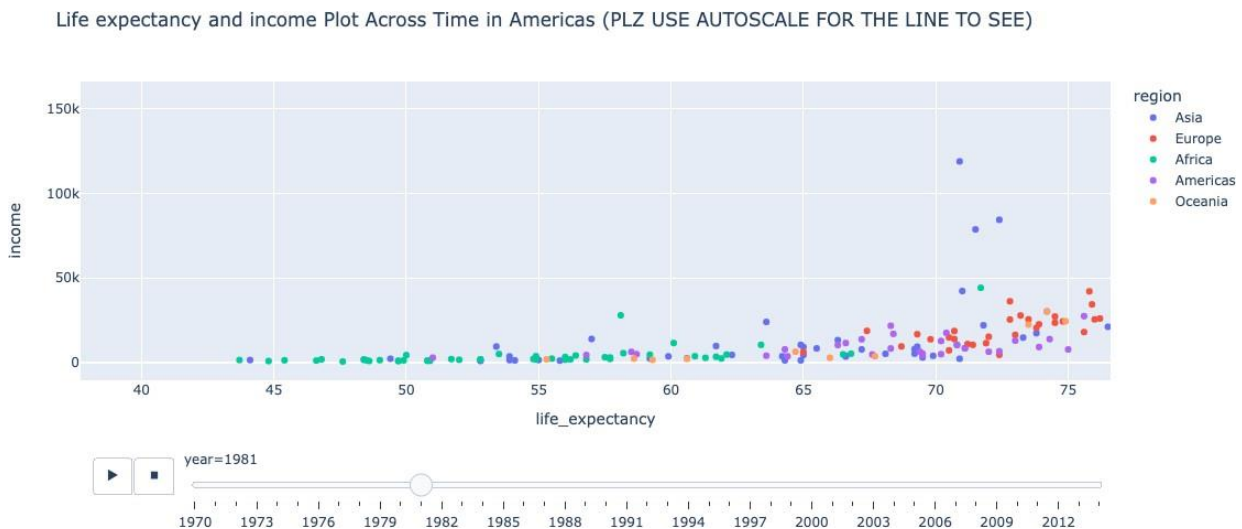


We see that, income seems very much aligned with life expectancy. As United states and Canada again score highest in Americans regions. On the other hand, Guatemala, Haiti & Honduras' income data had not changed likely life expectancy. Haiti's income value as well as Life Expectancy are lower than Guatemala's & Honduras

### 5.3.2 Let's plot Life Expectancy against Income to have a better understanding!

*## 2 Scatter plots to see the picture*

```
fig = px.scatter(data,x = "life_expectancy", y = "income", title = 'Life_
↳expectancy and income Plot Across Time in Americas (PLZ USE AUTOSCALE FOR_
↳THE LINE TO SEE)',
                color = "region",animation_frame="year", range_x=[0,1],_
↳range_y= [0,100])
fig.show()
```



We can see higher income is highly correlated with Life expectancy and gradually increase upto a certain level.

*## Lets see the relationship among different variables*

*# Pairplot to visualize relationships*


```
sns.pairplot(data)
```

[33]: *# Adjust the plot size*

```
plt.figure(figsize=(10, 15))
```

*# Show the plot*

```
plt.show()
```



<Figure size 1000x1500 with 0 Axes>

```
[34]: 4 # 3 Correlation Analysis
      ## Let's calculate the correlation matrix
      correlation_matrix = data.corr()
      print(correlation_matrix)
```

|                    | year      | population | life_expectancy | income    | children_per_woman | child_mortality | pop_density | co2_per_capita | years_in_school_men | years_in_school_women |
|--------------------|-----------|------------|-----------------|-----------|--------------------|-----------------|-------------|----------------|---------------------|-----------------------|
| year               | 1.000000  | 0.045262   | 0.318198        | 0.118619  | -0.383632          | -0.111739       | -0.810229   | -0.379525      | 0.534823            | 0.534823              |
| population         | 0.045262  | 1.000000   | 0.036585        | -0.025842 | -0.111739          | -0.111739       | -0.111739   | -0.111739      | 0.036585            | 0.036585              |
| life_expectancy    | 0.318198  | 0.036585   | 1.000000        | 0.534823  | -0.810229          | -0.379525       | -0.379525   | -0.379525      | 0.534823            | 0.534823              |
| income             | 0.118619  | -0.025842  | 0.534823        | 1.000000  | -0.379525          | -0.379525       | -0.379525   | -0.379525      | 0.534823            | 0.534823              |
| children_per_woman | -0.383632 | -0.111739  | -0.810229       | -0.379525 | 1.000000           | 0.036585        | 0.036585    | 0.036585       | -0.379525           | -0.379525             |

|                       |           |           |           |           |
|-----------------------|-----------|-----------|-----------|-----------|
| child_mortality       | -0.383741 | -0.037471 | -0.893179 | -0.443640 |
| pop_density           | 0.066457  | 0.014499  | 0.165861  | 0.152944  |
| co2_per_capita        | -0.013448 | 0.007397  | 0.428303  | 0.809062  |
| years_in_school_men   | 0.481321  | 0.052143  | 0.774083  | 0.448381  |
| years_in_school_women | 0.436158  | 0.010254  | 0.783480  | 0.462974  |

|                       | children_per_woman | child_mortality | pop_density | \ |
|-----------------------|--------------------|-----------------|-------------|---|
| year                  | -0.383632          | -0.383741       | 0.066457    |   |
| population            | -0.111739          | -0.037471       | 0.014499    |   |
| life_expectancy       | -0.810229          | -0.893179       | 0.165861    |   |
| income                | -0.379525          | -0.443640       | 0.152944    |   |
| children_per_woman    | 1.000000           | 0.839869        | -0.174382   |   |
| child_mortality       | 0.839869           | 1.000000        | -0.144161   |   |
| pop_density           | -0.174382          | -0.144161       | 1.000000    |   |
| co2_per_capita        | -0.362924          | -0.401001       | 0.082978    |   |
| years_in_school_men   | -0.835096          | -0.796552       | 0.112296    |   |
| years_in_school_women | -0.854199          | -0.804230       | 0.104472    |   |

|                       | co2_per_capita | years_in_school_men | \ |
|-----------------------|----------------|---------------------|---|
| year                  | -0.013448      | 0.481321            |   |
| population            | 0.007397       | 0.052143            |   |
| life_expectancy       | 0.428303       | 0.774083            |   |
| income                | 0.809062       | 0.448381            |   |
| children_per_woman    | -0.362924      | -0.835096           |   |
| child_mortality       | -0.401001      | -0.796552           |   |
| pop_density           | 0.082978       | 0.112296            |   |
| co2_per_capita        | 1.000000       | 0.387477            |   |
| years_in_school_men   | 0.387477       | 1.000000            |   |
| years_in_school_women | 0.412187       | 0.972736            |   |


|                       | years_in_school_women |
|-----------------------|-----------------------|
| year                  | 0.436158              |
| population            | 0.010254              |
| life_expectancy       | 0.783480              |
| income                | 0.462974              |
| children_per_woman    | -0.854199             |
| child_mortality       | -0.804230             |
| pop_density           | 0.104472              |
| co2_per_capita        | 0.412187              |
| years_in_school_men   | 0.972736              |
| years_in_school_women | 1.000000              |

<ipython-input-34-3d3b8c7d9b3a>:3: FutureWarning:

The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.



```
[35]: ## Visualizing the correlation matrix
plt.figure(figsize=(8, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Socio-Economic Indicators', fontsize=15)
plt.show()
```



## 6 Feature Selection

Feature selection is a crucial step in machine learning and data analysis. It involves choosing a subset of relevant features (variables) from your dataset to build a more effective model. Another step of the feature engineering process is to create classes of the categorical data and the resulting

labels are encoded as integers (0, 1, 2)

```
[36]: from sklearn.model_selection import train_test_split, GridSearchCV,
      ↪ cross_val_score
      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LinearRegression, LogisticRegression
      from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
      from sklearn.svm import SVR, SVC
      from sklearn.metrics import mean_squared_error, accuracy_score,
      ↪ classification_report

      # Considering features for
      features = ["year", "population", "life_expectancy", "children_per_woman",
      ↪ "child_mortality", "pop_density", "co2_per_capita", "years_in_school_men"]
      X = data[features]
      y_regression = data["income"] # Regression target
      y_classification = pd.qcut(data["life_expectancy"], q=3, labels=False) #
      ↪ Classification target, dividing life expectancy into 3 classes
```

##Data Preprocessing

```
[37]: # Split the data into train and test sets
      X_train, X_test, y_train_reg, y_test_reg = train_test_split(X, y_regression,
      ↪ test_size=0.2, random_state=42)
      X_train, X_test, y_train_cls, y_test_cls = train_test_split(X,
      ↪ y_classification, test_size=0.2, random_state=42)

      # Feature scaling
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

##Model Selection, Training, and Evaluation

Regression Models

```
[38]: # Linear Regression
      lr = LinearRegression()
      lr.fit(X_train_scaled, y_train_reg)
      lr_score = cross_val_score(lr, X_train_scaled, y_train_reg, cv=5,
      ↪ scoring='neg_mean_squared_error')
      print(f"Linear Regression RMSE: {np.sqrt(-lr_score.mean())}")

      # Random Forest Regressor
      rf_reg = RandomForestRegressor()
      rf_reg.fit(X_train_scaled, y_train_reg)
      rf_reg_score = cross_val_score(rf_reg, X_train_scaled, y_train_reg, cv=5,
      ↪ scoring='neg_mean_squared_error')
```



```

print(f"Random Forest Regressor RMSE: {np.sqrt(-rf_reg_score.mean())}")

# Support Vector Regressor
svr = SVR()
svr.fit(X_train_scaled, y_train_reg)
svr_score = cross_val_score(svr, X_train_scaled, y_train_reg, cv=5,
    ↳scoring='neg_mean_squared_error')
print(f"SVR RMSE: {np.sqrt(-svr_score.mean())}")

```

Linear Regression RMSE: 9146.103687168454  
 Random Forest Regressor RMSE: 3461.2778357076218  
 SVR RMSE: 18803.063264451783

---

## Classification Models

```

[ ]: #from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score, classification_report
#from sklearn.model_selection import cross_val_predict

# Assuming X_test_scaled and y_test_cls are your test sets

# Decision Tree Classifier
dt_cls = DecisionTreeClassifier()#
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
print(f"Decision Tree Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, dt_cls_pred)}")
print(f"Precision: {precision_score(y_test_cls, dt_cls_pred,
    ↳average='weighted'})")
print(f"Recall: {recall_score(y_test_cls, dt_cls_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, dt_cls_pred, average='weighted')}")
print(classification_report(y_test_cls, dt_cls_pred))

# K-Nearest Neighbors Classifier
knn_cls = KNeighborsClassifier()

```

```

knn_cls_pred = cross_val_predict(knn_cls, X_test_scaled, y_test_cls, cv=5)
print(f"K-Nearest Neighbors Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, knn_cls_pred)}")
print(f"Precision: {precision_score(y_test_cls, knn_cls_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, knn_cls_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, knn_cls_pred, average='weighted')}")
print(classification_report(y_test_cls, knn_cls_pred))

```

#### *# Logistic Regression Evaluation*

```
log_reg = LogisticRegression
```

```

knn_cls_pred = cross_val_predict(knn_cls, X_test_scaled, y_test_cls, cv=5)
print(f"K-Nearest Neighbors Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, knn_cls_pred)}")
print(f"Precision: {precision_score(y_test_cls, knn_cls_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, knn_cls_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, knn_cls_pred, average='weighted')}")
print(classification_report(y_test_cls, knn_cls_pred))

```

#### *# Logistic Regression Evaluation*

```

log_reg = LogisticRegression()#
log_reg_pred = cross_val_predict(log_reg, X_test_scaled, y_test_cls, cv=5)
print(f"Logistic Regression:")
print(f"Accuracy: {accuracy_score(y_test_cls, log_reg_pred)}")
print(f"Precision: {precision_score(y_test_cls, log_reg_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, log_reg_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, log_reg_pred, average='weighted')}")
print(classification_report(y_test_cls, log_reg_pred))

```

#### *# Random Forest Classifier Evaluation*

```

rf_cls = RandomForestClassifier()#
rf_cls_pred = cross_val_predict(rf_cls, X_test_scaled, y_test_cls, cv=5)
print(f"Random Forest Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, rf_cls_pred)}")
print(f"Precision: {precision_score(y_test_cls, rf_cls_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, rf_cls_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, rf_cls_pred, average='weighted')}")
print(classification_report(y_test_cls, rf_cls_pred))

```

#### *# Support Vector Classifier Evaluation*

```

svc = SVC()#
svc_pred = cross_val_predict(svc, X_test_scaled, y_test_cls, cv=5)
print(f"Support Vector Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, svc_pred)}")

```

```

print(f"Precision: {precision_score(y_test_cls, svc_pred, average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, svc_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, svc_pred, average='weighted')}")
print(classification_report(y_test_cls, svc_pred))

```

Decision Tree Classifier:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1-Score: 1.0

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 491     |
| 1            | 1.00      | 1.00   | 1.00     | 513     |
| 2            | 1.00      | 1.00   | 1.00     | 540     |
| accuracy     |           |        | 1.00     | 1544    |
| macro avg    | 1.00      | 1.00   | 1.00     | 1544    |
| weighted avg | 1.00      | 1.00   | 1.00     | 1544    |

K-Nearest Neighbors Classifier:

Accuracy: 0.9067357512953368

Precision: 0.9064837751919158

Recall: 0.9067357512953368

F1-Score: 0.9065946357646312

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.95      | 0.96   | 0.95     | 491     |
| 1            | 0.86      | 0.85   | 0.86     | 513     |
| 2            | 0.91      | 0.91   | 0.91     | 540     |
| accuracy     |           |        | 0.91     | 1544    |
| macro avg    | 0.91      | 0.91   | 0.91     | 1544    |
| weighted avg | 0.91      | 0.91   | 0.91     | 1544    |

Logistic Regression:

Accuracy: 0.9715025906735751

Precision: 0.9715436731234735

Recall: 0.9715025906735751

F1-Score: 0.971518945936874

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.98      | 0.98   | 0.98     | 491     |
| 1            | 0.96      | 0.96   | 0.96     | 513     |
| 2            | 0.98      | 0.98   | 0.98     | 540     |
| accuracy     |           |        | 0.97     | 1544    |
| macro avg    | 0.97      | 0.97   | 0.97     | 1544    |
| weighted avg | 0.97      | 0.97   | 0.97     | 1544    |

Accuracy: 0.9993523316062176

Precision: 0.9993535916614584

Recall: 0.9993523316062176

F1-Score: 0.9993523468023884

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 1.00      | 1.00   | 1.00     | 491     |
| 1 | 1.00      | 1.00   | 1.00     | 513     |

|              |      |      |      |      |
|--------------|------|------|------|------|
| 2            | 1.00 | 1.00 | 1.00 | 540  |
| accuracy     |      |      | 1.00 | 1544 |
| macro avg    | 1.00 | 1.00 | 1.00 | 1544 |
| weighted avg | 1.00 | 1.00 | 1.00 | 1544 |

Support Vector Classifier:

Accuracy: 0.9540155440414507

Precision: 0.9538671907249489

Recall: 0.9540155440414507

F1-Score: 0.9538865505825469

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.96      | 0.98   | 0.97     | 491     |
| 1            | 0.94      | 0.92   | 0.93     | 513     |
| 2            | 0.96      | 0.97   | 0.96     | 540     |
| accuracy     |           |        | 0.95     | 1544    |
| macro avg    | 0.95      | 0.95   | 0.95     | 1544    |
| weighted avg | 0.95      | 0.95   | 0.95     | 1544    |

### 3 To Check Effectiveness of the Models

For Classification Models we have already measured: “Accuracy”, “Precision”, “Recall” and “F1-Score”.

For Regression Models, we have measured: “RMSE (Root Mean Square Error)”

### 4 Now, to Check Efficiency of the Models

For Classification: We are measuring Cross-Validation time (which is a mix of “Training Time” and “Prediction/Testing Time” scores.)

We are using “Training Time” and “Prediction/Testing Time” which can help us understand the computational cost of your models.

```
[ ]: import time

## For Regression Models

# Linear Regression
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()
lr_pred = lr.predict(X_test_scaled)
```

```

end_pred = time.time()

print(f"Linear Regression Training Time: {end_train - start_train} seconds")
print(f"Linear Regression Prediction Time: {end_pred - start_pred} seconds")

# Random Forest Regressor
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()
lr_pred = lr.predict(X_test_scaled)
end_pred = time.time()

print(f"Random Forest Regressor Training Time: {end_train - start_train}_"
      ↪seconds")
print(f"Random Forest Regressor Prediction Time: {end_pred - start_pred}_"
      ↪seconds")

# Support Vector Regressor (SVR)
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()
lr_pred = lr.predict(X_test_scaled)
end_pred = time.time()

print(f"Support Vector Regressor (SVR) Training Time: {end_train - start_train}_"
      ↪seconds")
print(f"Support Vector Regressor (SVR) Prediction Time: {end_pred - start_pred}_"
      ↪seconds")

```

Linear Regression Training Time: 0.010734319686889648 seconds  
 Linear Regression Prediction Time: 0.0053863525390625 seconds  
 Random Forest Regressor Training Time: 0.009481191635131836 seconds  
 Random Forest Regressor Prediction Time: 0.0021293163299560547 seconds  
 Support Vector Regressor (SVR) Training Time: 0.004905223846435547 seconds  
 Support Vector Regressor (SVR) Prediction Time: 0.0004973411560058594 seconds

[ ]: *## For Classification Models*

```

# Decision Tree Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

```

```

print(f"Decision Tree Classifier Cross-Validation Time: {end - start} seconds")

# K-Nearest Neighbors Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"K-Nearest Neighbors Classifier Cross-Validation Time: {end - start}_
seconds")

# Logistic Regression
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Logistic Regression Cross-Validation Time: {end - start} seconds")

# Random Forest Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Random Forest Classifier Cross-Validation Time: {end - start} seconds")

# Support Vector Classifier (SVC)
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Support Vector Classifier (SVC) Cross-Validation Time: {end - start}_
seconds")

```

```

Decision Tree Classifier Cross-Validation Time: 0.04507708549499512 seconds
K-Nearest Neighbors Classifier Cross-Validation Time: 0.024494171142578125
seconds
Logistic Regression Cross-Validation Time: 0.02387261390686035 seconds
Random Forest Classifier Cross-Validation Time: 0.02438497543334961 seconds
Support Vector Classifier (SVC) Cross-Validation Time: 0.023949623107910156
seconds

```

## 5 To check Stability of the Models

We have already used “Cross-Validation Scores” to finalize among our models.

But Since our result shows that “Decision Tree” results the best scores in terms of “Accuracy”, “Precision”, “Recall” and “F1-Score” although it does not show efficiency in terms of “Cross-Validation Time” which is a mix of “Training Time” and “Prediction/Testing Time” scores.

#To ensure that our “Decision Tree” model’s high performance metrics are not a result of over-fitting, we employ more robust cross-validation technique to check the stability of “Decision Tree” accuracy.

```
[ ]: from sklearn.model_selection import StratifiedKFold, GridSearchCV
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import make_scorer, accuracy_score, precision_score,
      ↪ recall_score, f1_score
      from sklearn.model_selection import train_test_split

      # Assuming X and y_classification are your features and target variable for
      ↪ classification

      # Split the dataset into a training set and a test set to validate the model
      ↪ performance on unseen data
      X_train, X_test, y_train, y_test = train_test_split(X, y_classification,
      ↪ test_size=0.2, random_state=42)

      # Define a range of hyperparameters for the Decision Tree
      param_grid = {
          'criterion': ['gini', 'entropy'],
          'max_depth': [None, 10, 20, 30, 40, 50],
          'min_samples_split': [2, 5, 10],
          'min_samples_leaf': [1, 2, 4]
      }

      # Create a DecisionTreeClassifier instance
      dt = DecisionTreeClassifier(random_state=42)

      # Define the scoring function you want to optimize for, you can adjust these
      ↪ based on your requirements
      scoring = {'Accuracy': make_scorer(accuracy_score),
                  'Precision': make_scorer(precision_score, average='weighted'),
                  'Recall': make_scorer(recall_score, average='weighted'),
                  'F1': make_scorer(f1_score, average='weighted')}

      # Set up Stratified k-fold
      cv_strategy = StratifiedKFold(n_splits=10)

      # Set up GridSearchCV
      grid_search = GridSearchCV(dt, param_grid, scoring=scoring, refit='Accuracy',
      ↪ cv=cv_strategy, verbose=1, n_jobs=-1, return_train_score=True)

      # Fit the model
      grid_search.fit(X_train, y_train)

      # Best model result
```



```

best_model = grid_search.best_estimator_

# Validate on the test set
predictions = best_model.predict(X_test)
print("Test Set Evaluation:")
print(f"Accuracy: {accuracy_score(y_test, predictions)}")
print(f"Precision: {precision_score(y_test, predictions, average='weighted')}")
print(f"Recall: {recall_score(y_test, predictions, average='weighted')}")
print(f"F1-Score: {f1_score(y_test, predictions, average='weighted')}")

```

Fitting 10 folds for each of 108 candidates, totalling 1080 fits

Test Set Evaluation:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1-Score: 1.0

##Finalizing the Best Model

After running the above codes for analysing performance of models, we compare the RMSE for regression analysis models and accuracy, precision, recall as well as f1-score for classification models. Besides, accuracy score of “Decion tree” remains the same as before after checking the biaseness by overfitting, noisy or outlier-prone features of our dataset.

##The model with the lowest RMSE wih better efficiency (Random Forest Regressor) and highest accuracy, precision, recall as well as f1-score (Decision Tree Classifier), respectively, should be considered the best model for this particular dataset and research questions.

This approach provides a comprehensive answer to the posed research questions, utilizing machine learning to explore socio-economic indicators and predict future trends.

## 6 Key Notes and insights

The intial analysis has helped to draw the following conclusions about the data: \* The null values exist for a reason and imputing them using any other value would lead to wrong conclusions about the data set. \* Similarly, the presence of outliers is not due to error measurement. Removing these values would lead to wrong results and inferences about the data. \* The pair plot and correlation heat map help the most identifying the relationship of the different features.

asuring Cross-Validation time (which is a mix of “Training Time” and “Prediction/Testing Time” scores.)

We are using “Training Time” and “Prediction/Testing Time” which can help us understand the computational cost of your models.

```

import time

## For Regression Models

# Linear Regression
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()

```

```

end_pred = time.time()

print(f"Linear Regression Training Time: {end_train - start_train} seconds")
print(f"Linear Regression Prediction Time: {end_pred - start_pred} seconds")

# Random Forest Regressor
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()
lr_pred = lr.predict(X_test_scaled)
end_pred = time.time()

print(f"Random Forest Regressor Training Time: {end_train - start_train}_
↳seconds")
print(f"Random Forest Regressor Prediction Time: {end_pred - start_pred}_
↳seconds")

# Support Vector Regressor (SVR)
start_train = time.time()
lr.fit(X_train_scaled, y_train_reg)
end_train = time.time()

start_pred = time.time()
lr_pred = lr.predict(X_test_scaled)
end_pred = time.time()

print(f"Support Vector Regressor (SVR) Training Time: {end_train - start_train}_
↳seconds")
print(f"Support Vector Regressor (SVR) Prediction Time: {end_pred - start_pred}_
↳seconds")

```

Linear Regression Training Time: 0.010734319686889648 seconds  
 Linear Regression Prediction Time: 0.0053863525390625 seconds  
 Random Forest Regressor Training Time: 0.009481191635131836 seconds  
 Random Forest Regressor Prediction Time: 0.0021293163299560547 seconds  
 Support Vector Regressor (SVR) Training Time: 0.004905223846435547 seconds  
 Support Vector Regressor (SVR) Prediction Time: 0.0004973411560058594 seconds

[ ]: *## For Classification Models*

```

# Decision Tree Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

```

```

print(f"Decision Tree Classifier Cross-Validation Time: {end - start} seconds")

# K-Nearest Neighbors Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"K-Nearest Neighbors Classifier Cross-Validation Time: {end - start}_
seconds")

# Logistic Regression
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Logistic Regression Cross-Validation Time: {end - start} seconds")

# Random Forest Classifier
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Random Forest Classifier Cross-Validation Time: {end - start} seconds")

# Support Vector Classifier (SVC)
start = time.time()
dt_cls_pred = cross_val_predict(dt_cls, X_test_scaled, y_test_cls, cv=5)
end = time.time()

print(f"Support Vector Classifier (SVC) Cross-Validation Time: {end - start}_
seconds")

```

Decision Tree Classifier Cross-Validation Time: 0.04507708549499512 seconds

K-Nearest Neighbors Classifier Cross-Validation Time: 0.024494171142578125 seconds

Logistic Regression Cross-Validation Time: 0.02387261390686035 seconds

Random Forest Classifier Cross-Validation Time: 0.02438497543334961 seconds

Support Vector Classifier (SVC) Cross-Validation Time: 0.023949623107910156 seconds

## 7 To check Stability of the Models

We have already used “Cross-Validation Scores” to finalize among our models.

But Since our result shows that “Decision Tree” results the best scores in terms of “Accuracy”, “Precision”, “Recall” and “F1-Score” although it does not show efficiency in terms of “Cross-Validation Time” which is a mix of “Training Time” and “Prediction/Testing Time” scores.

#To ensure that our “Decision Tree” model’s high performance metrics are not a result of over-fitting, we employ more robust cross-validation technique to check the stability of “Decision Tree” accuracy.

```
[ ]: from sklearn.model_selection import StratifiedKFold, GridSearchCV
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import make_scorer, accuracy_score, precision_score,
      ↪ recall_score, f1_score
      from sklearn.model_selection import train_test_split

      # Assuming X and y_classification are your features and target variable for
      ↪ classification

      # Split the dataset into a training set and a test set to validate the model
      ↪ performance on unseen data
      X_train, X_test, y_train, y_test = train_test_split(X, y_classification,
      ↪ test_size=0.2, random_state=42)

      # Define a range of hyperparameters for the Decision Tree
      param_grid = {
          'criterion': ['gini', 'entropy'],
          'max_depth': [None, 10, 20, 30, 40, 50],
          'min_samples_split': [2, 5, 10],
          'min_samples_leaf': [1, 2, 4]
      }

      # Create a DecisionTreeClassifier instance
      dt = DecisionTreeClassifier(random_state=42)

      # Define the scoring function you want to optimize for, you can adjust these
      ↪ based on your requirements
      scoring = {'Accuracy': make_scorer(accuracy_score),
                  'Precision': make_scorer(precision_score, average='weighted'),
                  'Recall': make_scorer(recall_score, average='weighted'),
                  'F1': make_scorer(f1_score, average='weighted')}

      # Set up Stratified k-fold
      cv_strategy = StratifiedKFold(n_splits=10)

      # Set up GridSearchCV
      grid_search = GridSearchCV(dt, param_grid, scoring=scoring, refit='Accuracy',
      ↪ cv=cv_strategy, verbose=1, n_jobs=-1, return_train_score=True)

      # Fit the model
      grid_search.fit(X_train, y_train)

      # Best model result
```

```

best_model = grid_search.best_estimator_

# Validate on the test set
predictions = best_model.predict(X_test)
print("Test Set Evaluation:")
print(f"Accuracy: {accuracy_score(y_test, predictions)}")
print(f"Precision: {precision_score(y_test, predictions, average='weighted')}")
print(f"Recall: {recall_score(y_test, predictions, average='weighted')}")
print(f"F1-Score: {f1_score(y_test, predictions, average='weighted')}")

```

Fitting 10 folds for each of 108 candidates, totalling 1080 fits

Test Set Evaluation:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1-Score: 1.0

##Finalizing the Best Model

After running the above codes for analysing performance of models, we compare the RMSE for regression analysis models and accuracy, precision, recall as well as f1-score for classification models. Besides, accuracy score of “Decision tree” remains the same as before after checking the biasness by overfitting, noisy or outlier-prone features of our dataset.

##The model with the lowest RMSE with better efficiency (Random Forest Regressor) and highest accuracy, precision, recall as well as f1-score (Decision Tree Classifier), respectively, should be considered the best model for this particular dataset and research questions.

This approach provides a comprehensive answer to the posed research questions, utilizing machine learning to explore socio-economic indicators and predict future trends.

## 8 Key Notes and insights

The initial analysis has helped to draw the following conclusions about the data: \* The null values exist for a reason and imputing them using any other value would lead to wrong conclusions about the data set. \* Similarly, the presence of outliers is not due to error measurement. Removing these values would lead to wrong results and inferences about the data. \* The pair plot and correlation heat map help the most in identifying the relationship of the different features.