

# initial-results-and-the-code

March 30, 2024

This notebook presents the data cleaning, exploratory data analysis and feature engineering using data from [Gapminder](#) that will be used for linear, polynomial and logarithmic regression in another notebook.

## 1 Import packages and data

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns

# load the dataset (csv files)
data = pd.read_csv("https://raw.githubusercontent.com/BME1478H/Winter2022class/
↳master/data/world-data-gapminder.csv")
```

## 2 Description of data set and a summary of its attributes

This data set has been generated using data from [the Gapminder website](#), which focuses on gathering and sharing statistics and other information about social, economic and environmental development at local, national and global levels.

This particular data set describes the values of several parameters (see the list below) between 1800 and 2018 for a total of 192 countries. The parameters included in the data set and the column name of the dataframe are as follows:

- Country (country): Describes the country name
- Year (year): Describes the year to which the data belongs
- Population(population): Describes the number of population.
- Region (region): Describes the region to which the country belongs \*Sub Region (sub\_region): Describes the sub-region to which the country belongs.
- Income Group (income\_group): Classifies the people into income levels or classes.
- Life expectancy (life\_expectancy): Describes the life expectancy for a given country in a given year
- Income (income): Describes the Income per person in dollars for a given country in a given year
- Children per woman (children\_per-woman): Describes average number of children per woman
- Child Mortality (child-mortality): Describe number of child mortality in a country.
- Population Density (pop\_density): Describes the population density in a country.

- CO2 emissions per capita (co2\_per\_capita): Describes the CO2 emissions in tonnes per person for a given country in a given year
- Years in school for Men (years\_in\_school\_men): Describe the number of years on average a men spent in school
- Years in school for Women (years\_in\_school\_women): Describe the number of years on average an women spent in school

Let's have a look at the data.

Let's check the data types and the number of samples for each column:

```
[2]: #Dataset chekup
data.sample(n=10)
```

```
[2]:      country  year  population  region \
17124   Jamaica 1842    400000  Americas
29762   Senegal 1997   9200000   Africa
219     Albania 1800    410000  Europe
8220   Costa Rica 1917    400000  Americas
24398    Nepal 1889   5240000   Asia
34974    Togo 1953   1450000   Africa
1961    Austria 2009   8370000  Europe
1638   Australia 1905   4020000  Oceania
14796    Haiti 1923   2230000  Americas
27204   Paraguay 1848    318000  Americas

      sub_region  income_group  life_expectancy  income \
17124 Latin America and the Caribbean  Upper middle      34.2      977
29762          Sub-Saharan Africa      Low      58.8     1760
219          Southern Europe  Upper middle      35.4      667
8220 Latin America and the Caribbean  Upper middle      35.6     2300
24398          Southern Asia      Low      33.8      758
34974          Sub-Saharan Africa      Low      39.0     1080
1961          Western Europe      High      80.3    42500
1638  Australia and New Zealand      High      52.7     7140
14796 Latin America and the Caribbean      Low      28.5     1810
27204 Latin America and the Caribbean  Upper middle      35.5     1190

      children_per_woman  child_mortality  pop_density  co2_per_capita \
17124                5.13            390.0         NaN           NaN
29762                5.77            143.0         47.8       0.3550
219                4.60            375.0         NaN           NaN
8220                6.71            366.0         NaN           NaN
24398                6.15            407.0         NaN           NaN
34974                6.33            308.0         26.7       0.0202
1961                1.39              4.5        102.0       7.4900
1638                3.51            117.0         NaN       3.0000
14796                6.31            462.0         NaN           NaN
```

27204	6.49	374.0	NaN	NaN
-------	------	-------	-----	-----

	years_in_school_men	years_in_school_women
17124	NaN	NaN
29762	3.34	1.82
219	NaN	NaN
8220	NaN	NaN
24398	NaN	NaN
34974	NaN	NaN
1961	12.40	12.40
1638	NaN	NaN
14796	NaN	NaN
27204	NaN	NaN

```
[3]: # Display the first few rows of the dataset to understand its structure
print(data.head())
```

	country	year	population	region	sub_region	income_group	\
0	Afghanistan	1800	3280000	Asia	Southern Asia	Low	
1	Afghanistan	1801	3280000	Asia	Southern Asia	Low	
2	Afghanistan	1802	3280000	Asia	Southern Asia	Low	
3	Afghanistan	1803	3280000	Asia	Southern Asia	Low	
4	Afghanistan	1804	3280000	Asia	Southern Asia	Low	

	life_expectancy	income	children_per_woman	child_mortality	pop_density	\
0	28.2	603	7.0	469.0	NaN	
1	28.2	603	7.0	469.0	NaN	
2	28.2	603	7.0	469.0	NaN	
3	28.2	603	7.0	469.0	NaN	
4	28.2	603	7.0	469.0	NaN	

	co2_per_capita	years_in_school_men	years_in_school_women
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

```
[4]: # Display data types
print(data.dtypes)
```

country	object
year	int64
population	int64
region	object
sub_region	object
income_group	object
life_expectancy	float64

```

income                int64
children_per_woman    float64
child_mortality        float64
pop_density            float64
co2_per_capita         float64
years_in_school_men    float64
years_in_school_women  float64
dtype: object

```

[5]: *# List of columns for reference in the analyses below*  
`print(data.columns)`

```

Index(['country', 'year', 'population', 'region', 'sub_region', 'income_group',
      'life_expectancy', 'income', 'children_per_woman', 'child_mortality',
      'pop_density', 'co2_per_capita', 'years_in_school_men',
      'years_in_school_women'],
      dtype='object')

```

[6]: *# For categorical attributes, let's look at the frequency of countries or other\_*  
*↪categorical fields*  
`print(data["country"].value_counts())` *# Assuming 'country' is one of the\_*  
*↪columns*

```

United States    438
Afghanistan      219
Panama           219
New Zealand      219
Nicaragua        219
...
Greece           219
Grenada          219
Guatemala        219
Guinea           219
Zimbabwe         219
Name: country, Length: 178, dtype: int64

```

[7]: *# Basic information about the dataset*  
`print(data.info())`  
*#data.info()*

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39201 entries, 0 to 39200
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country                39201 non-null  object
1   year                   39201 non-null  int64
2   population              39201 non-null  int64

```

3	region	39201	non-null	object
4	sub_region	39201	non-null	object
5	income_group	39201	non-null	object
6	life_expectancy	39201	non-null	float64
7	income	39201	non-null	int64
8	children_per_woman	39201	non-null	float64
9	child_mortality	39199	non-null	float64
10	pop_density	12351	non-null	float64
11	co2_per_capita	16500	non-null	float64
12	years_in_school_men	8234	non-null	float64
13	years_in_school_women	8234	non-null	float64

dtypes: float64(7), int64(3), object(4)

memory usage: 4.2+ MB

None

### 3 Initial plan for data exploration

The initial plan for data exploration is as follows:

\*Data cleaning: The data cleaning process is split into following two main actions . . . Null values: Understand the reason why there are null values to find out the best way to deal with them.

Outliers: Similarly to the null values, the first step is to understand the presence of outliers as well as to find out if removing them is a good idea or it's actually valuable data.

\*Feature extraction: Understand the relationship between the different features, perform transformation to help improve those relationships and perform Principal Component Analysis to understand how some of the features explain the HDI index variance

\*Feature extraction: Understand the relationship between the different features, perform transformation to help improve those relationships and perform Principal Component Analysis to understand how some of the features explain the HDI index variance

\*Hypothesis testing: Formulate three hypotheses and test one of them

Next, Let's find out how many null values there are:

```
[8]: # To find out the number of Null values
data.isnull().sum()
```

```
[8]: country          0
     year             0
     population       0
     region           0
     sub_region       0
     income_group     0
     life_expectancy  0
     income           0
     children_per_woman 0
     child_mortality  2
```

```

pop_density          26850
co2_per_capita       22701
years_in_school_men  30967
years_in_school_women 30967
dtype: int64

```

## 4 Data cleaning

### 4.1 Null values

Data cleaning is always tricky and this dataset isn't an exception. In fact, it's even trickier because the number of samples, 39202, is spread accross 218 different years. This means that there are 218 different sub-datasets with around 205 elements each. Therefore, the sample size is relatively small and the impact of each value in the final metrics is greater. One needs to be careful and undrestand the reason behind the NaN values before replacing or deleting them.

Prior to the analysis, I considered the following techniques to deal with null values:

- Remove them, provided that we don't lose a lot of data
- Impute data using the one of the following options:
  - Mean value
  - Most common value
  - Interpolation provided that there is enough data to make a prediction

However, let's starts by checking which are the null values and try to understand why there are missing values.

```
[9]: data[data.isna().any(axis=1)]
```

```

[9]:      country  year  population  region  sub_region income_group \
0    Afghanistan  1800    3280000  Asia    Southern Asia    Low
1    Afghanistan  1801    3280000  Asia    Southern Asia    Low
2    Afghanistan  1802    3280000  Asia    Southern Asia    Low
3    Afghanistan  1803    3280000  Asia    Southern Asia    Low
4    Afghanistan  1804    3280000  Asia    Southern Asia    Low
...
39151  Zimbabwe  1969    5010000  Africa  Sub-Saharan Africa    Low
39197  Zimbabwe  2015   15800000  Africa  Sub-Saharan Africa    Low
39198  Zimbabwe  2016   16200000  Africa  Sub-Saharan Africa    Low
39199  Zimbabwe  2017   16500000  Africa  Sub-Saharan Africa    Low
39200  Zimbabwe  2018   16900000  Africa  Sub-Saharan Africa    Low

      life_expectancy  income  children_per_woman  child_mortality \
0                28.2    603                7.00            469.0
1                28.2    603                7.00            469.0
2                28.2    603                7.00            469.0
3                28.2    603                7.00            469.0
4                28.2    603                7.00            469.0

```

...	...	...	...	...
39151	57.2	2160	7.42	115.0
39197	58.3	1890	3.84	59.9
39198	59.3	1860	3.76	56.4
39199	59.8	1910	3.68	56.8
39200	60.2	1950	3.61	55.5

	pop_density	co2_per_capita	years_in_school_men	years_in_school_women
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

...	...	...	...	...
39151	12.9	1.35	NaN	NaN
39197	40.8	NaN	11.1	10.2
39198	41.7	NaN	NaN	NaN
39199	42.7	NaN	NaN	NaN
39200	43.7	NaN	NaN	NaN

[31484 rows x 14 columns]

Let's have a look at the list of unique countries and years whose rows include null values for the pop\_densityColumn:

```
[10]: print(set(data[data["pop_density"].isna()]["country"]))
print(set(data[data["pop_density"].isna()]["year"]))
```

```
{'Gambia', 'Montenegro', 'Swaziland', 'Iraq', 'Peru', 'Tunisia', 'Congo, Dem. Rep.', 'Zimbabwe', 'Uganda', 'Chile', 'Czech Republic', 'Oman', 'Malaysia', 'Lithuania', 'Australia', 'Mauritania', 'Estonia', 'Poland', 'Netherlands', 'Bolivia', 'Sri Lanka', 'Benin', 'Italy', 'Niger', 'Slovenia', 'Cambodia', 'Afghanistan', 'Solomon Islands', 'Spain', 'Malawi', 'Gabon', 'Ghana', 'Canada', 'Denmark', 'Latvia', 'Norway', 'France', 'Equatorial Guinea', 'Jamaica', 'Guinea-Bissau', 'Senegal', 'Moldova', 'Burundi', 'Bahamas', 'Burkina Faso', 'Argentina', 'Liberia', 'South Sudan', 'Venezuela', 'Kiribati', 'Mexico', 'Switzerland', 'El Salvador', 'New Zealand', 'Sweden', 'Kazakhstan', 'South Africa', 'Samoa', 'Croatia', 'Djibouti', 'Trinidad and Tobago', 'Brazil', 'Thailand', 'Tonga', 'Kenya', 'Slovak Republic', 'Ireland', 'Central African Republic', 'Bhutan', 'Lebanon', 'Panama', 'Cuba', 'Serbia', 'Uzbekistan', 'Papua New Guinea', 'Greece', 'Bulgaria', 'Bangladesh', 'Saudi Arabia', 'Lao', 'Mongolia', 'Ecuador', 'Cyprus', 'United Arab Emirates', 'Somalia', 'Vietnam', 'Seychelles', 'Paraguay', 'Palestine', 'Romania', 'Libya', 'Botswana', 'Algeria', 'Germany', 'Zambia', 'Tajikistan', 'Azerbaijan', 'Georgia', 'Ethiopia', 'Macedonia, FYR', 'United Kingdom', 'Uruguay', 'Suriname', 'Costa Rica', 'Egypt', 'Guinea', 'Pakistan', 'Hungary', 'Rwanda', 'Mozambique', 'Mauritius', 'Kuwait', 'Togo', 'Kyrgyz Republic', 'Syria', 'Iceland', 'Chad',
```

'Barbados', 'Qatar', 'Albania', 'Nepal', 'Namibia', 'Dominican Republic',  
 'Guatemala', 'Cote d'Ivoire', 'India', 'Bahrain', 'Belize', 'Nicaragua',  
 'Comoros', 'Finland', 'Malta', 'United States', 'Grenada', 'Israel', 'Nigeria',  
 'China', 'Armenia', 'Luxembourg', 'Japan', 'Tanzania', 'Haiti', 'Morocco',  
 'Bosnia and Herzegovina', 'Congo, Rep.', 'Mali', 'Sierra Leone', 'South Korea',  
 'Singapore', 'Fiji', 'Myanmar', 'Turkey', 'Jordan', 'North Korea', 'Sudan',  
 'Honduras', 'Belarus', 'Austria', 'Yemen', 'Maldives', 'Eritrea', 'Angola',  
 'Colombia', 'Iran', 'Belgium', 'Indonesia', 'Portugal', 'Russia', 'Timor-Leste',  
 'Turkmenistan', 'Cameroon', 'Madagascar', 'Guyana', 'Lesotho', 'Vanuatu',  
 'Antigua and Barbuda', 'Philippines', 'Ukraine'}

{1800, 1801, 1802, 1803, 1804, 1805, 1806, 1807, 1808, 1809, 1810, 1811, 1812,  
 1813, 1814, 1815, 1816, 1817, 1818, 1819, 1820, 1821, 1822, 1823, 1824, 1825,  
 1826, 1827, 1828, 1829, 1830, 1831, 1832, 1833, 1834, 1835, 1836, 1837, 1838,  
 1839, 1840, 1841, 1842, 1843, 1844, 1845, 1846, 1847, 1848, 1849, 1850, 1851,  
 1852, 1853, 1854, 1855, 1856, 1857, 1858, 1859, 1860, 1861, 1862, 1863, 1864,  
 1865, 1866, 1867, 1868, 1869, 1870, 1871, 1872, 1873, 1874, 1875, 1876, 1877,  
 1878, 1879, 1880, 1881, 1882, 1883, 1884, 1885, 1886, 1887, 1888, 1889, 1890,  
 1891, 1892, 1893, 1894, 1895, 1896, 1897, 1898, 1899, 1900, 1901, 1902, 1903,  
 1904, 1905, 1906, 1907, 1908, 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916,  
 1917, 1918, 1919, 1920, 1921, 1922, 1923, 1924, 1925, 1926, 1927, 1928, 1929,  
 1930, 1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942,  
 1943, 1944, 1945, 1946, 1947, 1948, 1949}

It makes sense to see some of these countries and years as some of them declared their independence recently, others underwent foreign military occupations or wars.

Next, Let's follow a similar approach for the 'co2\_per\_capita' column:

```
[11]: data[data["co2_per_capita"].isna()].sample(n=10)
```

```
[11]:      country  year  population  region  sub_region \
31094  Slovenia  2015    2070000  Europe  Southern Europe
38406  Vietnam  1881   11600000  Asia    South-eastern Asia
19826  Liberia  1916    702000  Africa  Sub-Saharan Africa
19511  Lesotho  1820    282000  Africa  Sub-Saharan Africa
32916  Suriname  1866    114000  Americas Latin America and the Caribbean
17626  Jordan   1906    342000  Asia    Western Asia
13387  Greece   1828   2490000  Europe  Southern Europe
38596  Yemen    1852   2750000  Asia    Western Asia
16250  Iraq     1844   1310000  Asia    Western Asia
15019  Honduras 1927    942000  Americas Latin America and the Caribbean
```

```
      income_group  life_expectancy  income  children_per_woman \
31094           High             80.7   29100             1.61
38406  Lower middle             31.7     917             4.70
19826           Low             34.1    1030             6.19
19511  Lower middle             32.8     398             5.84
32916  Upper middle             32.9    2240             6.58
```



17626	Upper middle	32.1	1590	6.97
13387	High	36.6	1520	6.03
38596	Low	23.4	1020	6.88
16250	Upper middle	31.2	1070	7.13
15019	Lower middle	35.7	2650	6.35

	child_mortality	pop_density	co2_per_capita	years_in_school_men	\
31094	2.4	103.0	NaN	13.4	
38406	417.0	NaN	NaN	NaN	
19826	416.0	NaN	NaN	NaN	
19511	407.0	NaN	NaN	NaN	
32916	406.0	NaN	NaN	NaN	
17626	417.0	NaN	NaN	NaN	
13387	361.0	NaN	NaN	NaN	
38596	540.0	NaN	NaN	NaN	
16250	428.0	NaN	NaN	NaN	
15019	371.0	NaN	NaN	NaN	

	years_in_school_women
31094	14.2
38406	NaN
19826	NaN
19511	NaN
32916	NaN
17626	NaN
13387	NaN
38596	NaN
16250	NaN
15019	NaN

```
[12]: print(set(data[data["co2_per_capita"].isna()]["country"]))
      print(set(data[data["co2_per_capita"].isna()]["year"]))
```

```
{'Gambia', 'Montenegro', 'Swaziland', 'Iraq', 'Peru', 'Tunisia', 'Congo, Dem. Rep.', 'Zimbabwe', 'Uganda', 'Chile', 'Czech Republic', 'Oman', 'Malaysia', 'Lithuania', 'Australia', 'Mauritania', 'Estonia', 'Poland', 'Netherlands', 'Bolivia', 'Sri Lanka', 'Benin', 'Italy', 'Niger', 'Slovenia', 'Cambodia', 'Afghanistan', 'Solomon Islands', 'Spain', 'Malawi', 'Gabon', 'Ghana', 'Canada', 'Denmark', 'Latvia', 'Norway', 'France', 'Equatorial Guinea', 'Jamaica', 'Guinea-Bissau', 'Senegal', 'Moldova', 'Burundi', 'Bahamas', 'Burkina Faso', 'Argentina', 'Liberia', 'South Sudan', 'Venezuela', 'Kiribati', 'Mexico', 'Switzerland', 'El Salvador', 'New Zealand', 'Sweden', 'Kazakhstan', 'South Africa', 'Samoa', 'Croatia', 'Djibouti', 'Trinidad and Tobago', 'Brazil', 'Thailand', 'Tonga', 'Kenya', 'Slovak Republic', 'Ireland', 'Central African Republic', 'Bhutan', 'Lebanon', 'Panama', 'Cuba', 'Serbia', 'Uzbekistan', 'Papua New Guinea', 'Greece', 'Bulgaria', 'Bangladesh', 'Saudi Arabia', 'Lao'}
```

'Mongolia', 'Ecuador', 'Cyprus', 'United Arab Emirates', 'Somalia', 'Vietnam', 'Seychelles', 'Paraguay', 'Palestine', 'Romania', 'Libya', 'Botswana', 'Algeria', 'Germany', 'Zambia', 'Tajikistan', 'Azerbaijan', 'Georgia', 'Ethiopia', 'Macedonia, FYR', 'United Kingdom', 'Uruguay', 'Suriname', 'Costa Rica', 'Egypt', 'Guinea', 'Pakistan', 'Hungary', 'Rwanda', 'Mozambique', 'Mauritius', 'Kuwait', 'Togo', 'Kyrgyz Republic', 'Syria', 'Iceland', 'Chad', 'Barbados', 'Qatar', 'Albania', 'Nepal', 'Namibia', 'Dominican Republic', 'Guatemala', 'Cote d'Ivoire', 'India', 'Bahrain', 'Belize', 'Nicaragua', 'Comoros', 'Finland', 'Malta', 'United States', 'Grenada', 'Israel', 'Nigeria', 'China', 'Armenia', 'Luxembourg', 'Japan', 'Tanzania', 'Haiti', 'Morocco', 'Bosnia and Herzegovina', 'Congo, Rep.', 'Mali', 'Sierra Leone', 'South Korea', 'Singapore', 'Fiji', 'Myanmar', 'Turkey', 'Jordan', 'North Korea', 'Sudan', 'Honduras', 'Belarus', 'Austria', 'Yemen', 'Maldives', 'Eritrea', 'Angola', 'Colombia', 'Iran', 'Belgium', 'Indonesia', 'Portugal', 'Russia', 'Timor-Leste', 'Turkmenistan', 'Cameroon', 'Madagascar', 'Guyana', 'Lesotho', 'Vanuatu', 'Antigua and Barbuda', 'Philippines', 'Ukraine'}

{1800, 1801, 1802, 1803, 1804, 1805, 1806, 1807, 1808, 1809, 1810, 1811, 1812, 1813, 1814, 1815, 1816, 1817, 1818, 1819, 1820, 1821, 1822, 1823, 1824, 1825, 1826, 1827, 1828, 1829, 1830, 1831, 1832, 1833, 1834, 1835, 1836, 1837, 1838, 1839, 1840, 1841, 1842, 1843, 1844, 1845, 1846, 1847, 1848, 1849, 1850, 1851, 1852, 1853, 1854, 1855, 1856, 1857, 1858, 1859, 1860, 1861, 1862, 1863, 1864, 1865, 1866, 1867, 1868, 1869, 1870, 1871, 1872, 1873, 1874, 1875, 1876, 1877, 1878, 1879, 1880, 1881, 1882, 1883, 1884, 1885, 1886, 1887, 1888, 1889, 1890, 1891, 1892, 1893, 1894, 1895, 1896, 1897, 1898, 1899, 1900, 1901, 1902, 1903, 1904, 1905, 1906, 1907, 1908, 1909, 1910, 1911, 1912, 1913, 1914, 1915, 1916, 1917, 1918, 1919, 1920, 1921, 1922, 1923, 1924, 1925, 1926, 1927, 1928, 1929, 1930, 1931, 1932, 1933, 1934, 1935, 1936, 1937, 1938, 1939, 1940, 1941, 1942, 1943, 1944, 1945, 1946, 1947, 1948, 1949, 1950, 1951, 1952, 1953, 1954, 1955, 1956, 1957, 1958, 1959, 1960, 1961, 1962, 1963, 1964, 1965, 1966, 1967, 1968, 1969, 1970, 1971, 1972, 1973, 1974, 1975, 1976, 1977, 1978, 1979, 1980, 1981, 1982, 1983, 1984, 1985, 1986, 1987, 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2015, 2016, 2017, 2018}

We can also see countries that have undergone difficult periods. Given the complexity of these events, it will be difficult to make predictions or find the right value to replace for all countries, i.e. the mean value might be good for some but it might lead to uncertainty in others. Furthermore, we can see that there is a period (2013-2015) in which there are no null values. Therefore, it seems like the best approach is to **remove the null values**.

```
[13]: data = data.dropna().reset_index(drop=True)
      data.isnull().sum()
```

```
[13]: country      0
      year         0
      population   0
      region       0
```

```

sub_region      0
income_group    0
life_expectancy 0
income          0
children_per_woman 0
child_mortality 0
pop_density     0
co2_per_capita  0
years_in_school_men 0
years_in_school_women 0
dtype: int64

```

## 4.2 Select a subset

Given the range of the data (1800-2018), drawing conclusions based on the entire data set is difficult and counter productive. We can expect similar relationship for adjacent years, such as 1995-1998, 1998-2000, 2005-2007 and so on, but most likely those relationship have changed between 1800 and 2018. Therefore, working with the entire data set might lead to not seeing the full picture.

This being said, we can select a subset of the data and in the data cleaning section we identified a subset that included no null values and it isn't affected by the loss of information.

```

[14]: # select the subset belonging to the year range 2013-2015
data_period = data[
    (data["year"] == 2013) | (data["year"] == 2014) | (data["year"] == 2015)
]
data_period.head()

```

```

[14]:      country  year  population  region  sub_region  income_group \
43  Afghanistan  2013   31700000  Asia    Southern Asia      Low
44  Afghanistan  2014   32800000  Asia    Southern Asia      Low
88    Albania   2013    2920000  Europe  Southern Europe  Upper middle
89    Albania   2014    2920000  Europe  Southern Europe  Upper middle
133  Algeria    2013   38300000  Africa  Northern Africa  Upper middle

```

```

      life_expectancy  income  children_per_woman  child_mortality \
43              57.7     1810              5.17             79.3
44              57.8     1780              4.98             76.1
88              77.2    10500              1.70             14.9
89              77.4    10700              1.71             14.4
133             77.0    13300              2.92             25.8

```

```

      pop_density  co2_per_capita  years_in_school_men  years_in_school_women
43           48.6           0.316              3.94              0.92
44           50.2           0.299              4.04              0.95
88          107.0           1.730             11.70             11.90
89          107.0           1.960             11.80             12.10

```

133            16.1            3.510            8.24            7.42

[15]: data\_period.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 358 entries, 43 to 7716

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	country	358 non-null	object
1	year	358 non-null	int64
2	population	358 non-null	int64
3	region	358 non-null	object
4	sub_region	358 non-null	object
5	income_group	358 non-null	object
6	life_expectancy	358 non-null	float64
7	income	358 non-null	int64
8	children_per_woman	358 non-null	float64
9	child_mortality	358 non-null	float64
10	pop_density	358 non-null	float64
11	co2_per_capita	358 non-null	float64
12	years_in_school_men	358 non-null	float64
13	years_in_school_women	358 non-null	float64

dtypes: float64(7), int64(3), object(4)

memory usage: 42.0+ KB

[16]: data\_period.groupby("region").describe()

```
[16]:
```

	year	count	mean	std	min	25%	50%	75%	max
region									
Africa	104.0	2013.5	0.502421	2013.0	2013.0	2013.5	2014.0	2014.0	
Americas	62.0	2013.5	0.504082	2013.0	2013.0	2013.5	2014.0	2014.0	
Asia	94.0	2013.5	0.502681	2013.0	2013.0	2013.5	2014.0	2014.0	
Europe	78.0	2013.5	0.503236	2013.0	2013.0	2013.5	2014.0	2014.0	
Oceania	20.0	2013.5	0.512989	2013.0	2013.0	2013.5	2014.0	2014.0	

```

      population
      count      mean
region
Africa      104.0  2.205035e+07
Americas     62.0  3.124632e+07
Asia         94.0  9.190187e+07
Europe       78.0  1.896846e+07
Oceania      20.0  3.546855e+07

      years_in_school_men
      75%      max
region
Africa      8.9275  11.3
Americas    11.8750  15.3
Asia        12.4000  15.0
Europe      13.9000  14.8
Oceania     13.9250  14.5

      years_in_school_women

```

	count	mean	std	min	25%	50%
region						
Africa	104.0	5.646827	2.886351	1.12	3.5250	5.38
Americas	62.0	11.053226	2.121167	6.35	9.7200	10.85
Asia	94.0	9.859362	3.658881	0.92	7.5700	10.60
Europe	78.0	13.733333	1.039938	11.00	13.3250	14.00
Oceania	20.0	11.268000	3.093146	5.75	8.7375	11.90

	75%	max
region		
Africa	7.795	11.4
Americas	12.700	15.5
Asia	12.800	15.6
Europe	14.400	15.3
Oceania	14.325	14.9

[5 rows x 80 columns]

## 5 Outliers

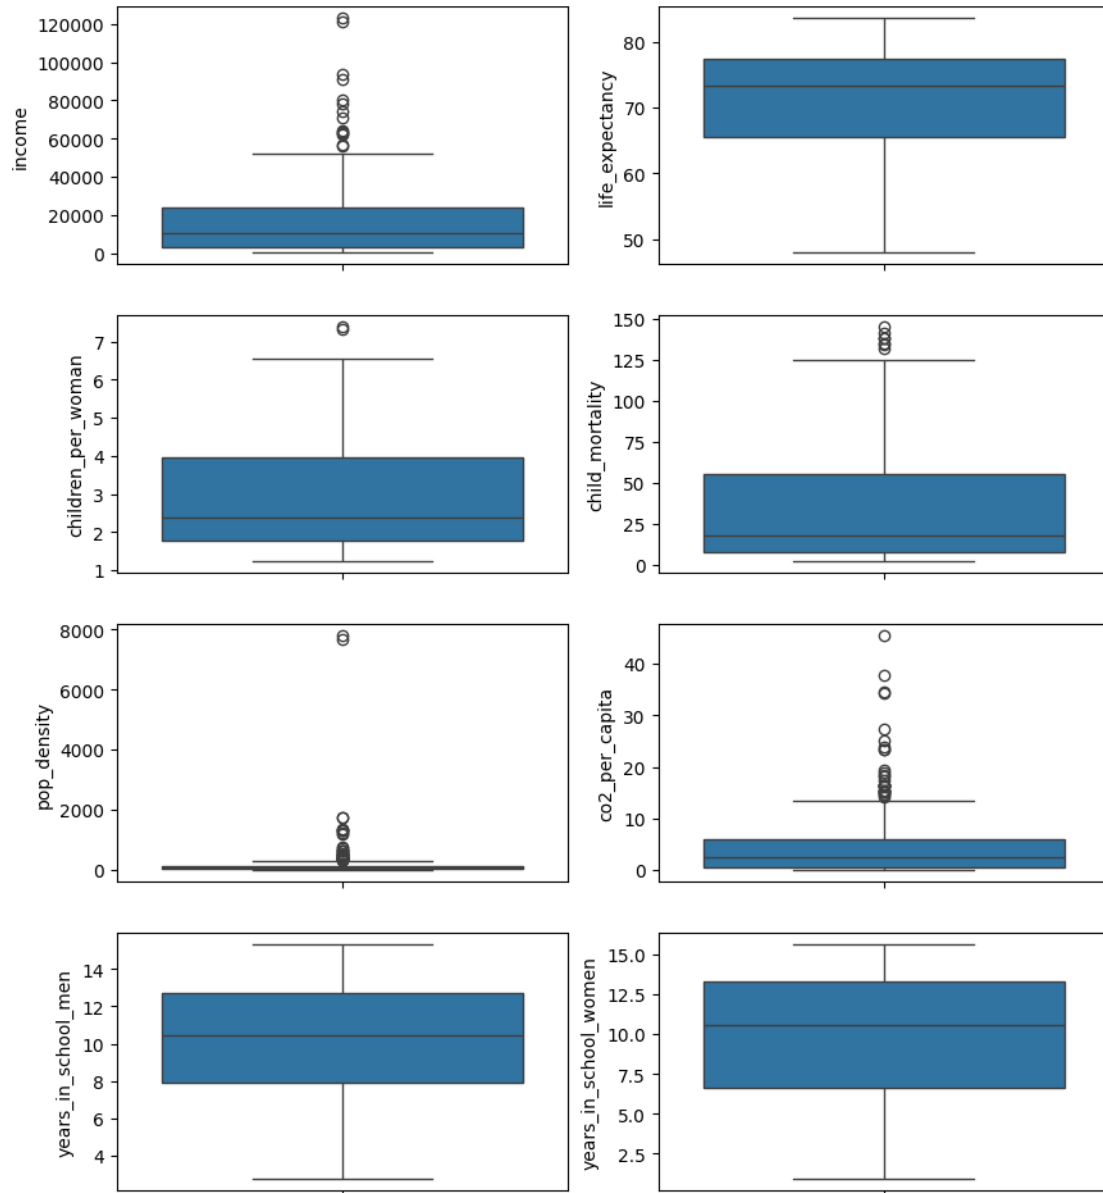
The presence of outliers can easily be identified by using box plots. However, one has to be careful as a boxplot of the entire data frame will show many outliers. Let's create the box plots

```
[17]: # Create a list using the features that contain numerical data
column_list_plot = [
    "income",
    "life_expectancy",
    "children_per_woman",
    "child_mortality",
    "pop_density",
    "co2_per_capita",
    "years_in_school_men",
    "years_in_school_women",
]

# Create a 4x2 figure with 8 subplots, where 8 of them will be used
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(10, 12))

# Flatten the axes array to iterate over it
axes = axes.flatten()

# Use a for loop to create the subplots
for index, col_name in enumerate(column_list_plot):
    sns.boxplot(ax=axes[index], y=col_name, data=data_period)
```



As expected, the box plots for four of the features (income, child\_mortality, pop\_density and co2\_emissions) show many outliers, but this isn't the full picture. An accurate analysis of the outliers requires to plot the data split into continents.

[18]: *# Create a list of colors for different regions*

```
region_colors = {
    "Africa": "green",
    "Americas": "yellow",
    "Asia": "blue",
    "Europe": "purple",
    "Oceania": "red",
```

```

}

# Create a 4x2 figure with 8 subplots, Where 8 of them will be used
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(12, 15))

# Flatten the axes array to iterate over it
axes = axes.flatten()

# Use a for loop to create the subplots
for index, col_name in enumerate(column_list_plot):
    sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
               palette=region_colors)

# Add a title to the entire figure
fig.suptitle("Box Plots of Numerical Features by Region", fontsize=20)

```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
palette=region_colors)
```

<ipython-input-18-28821a18bb68>:18: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,
```

```
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

```
<ipython-input-18-28821a18bb68>:18: FutureWarning:
```

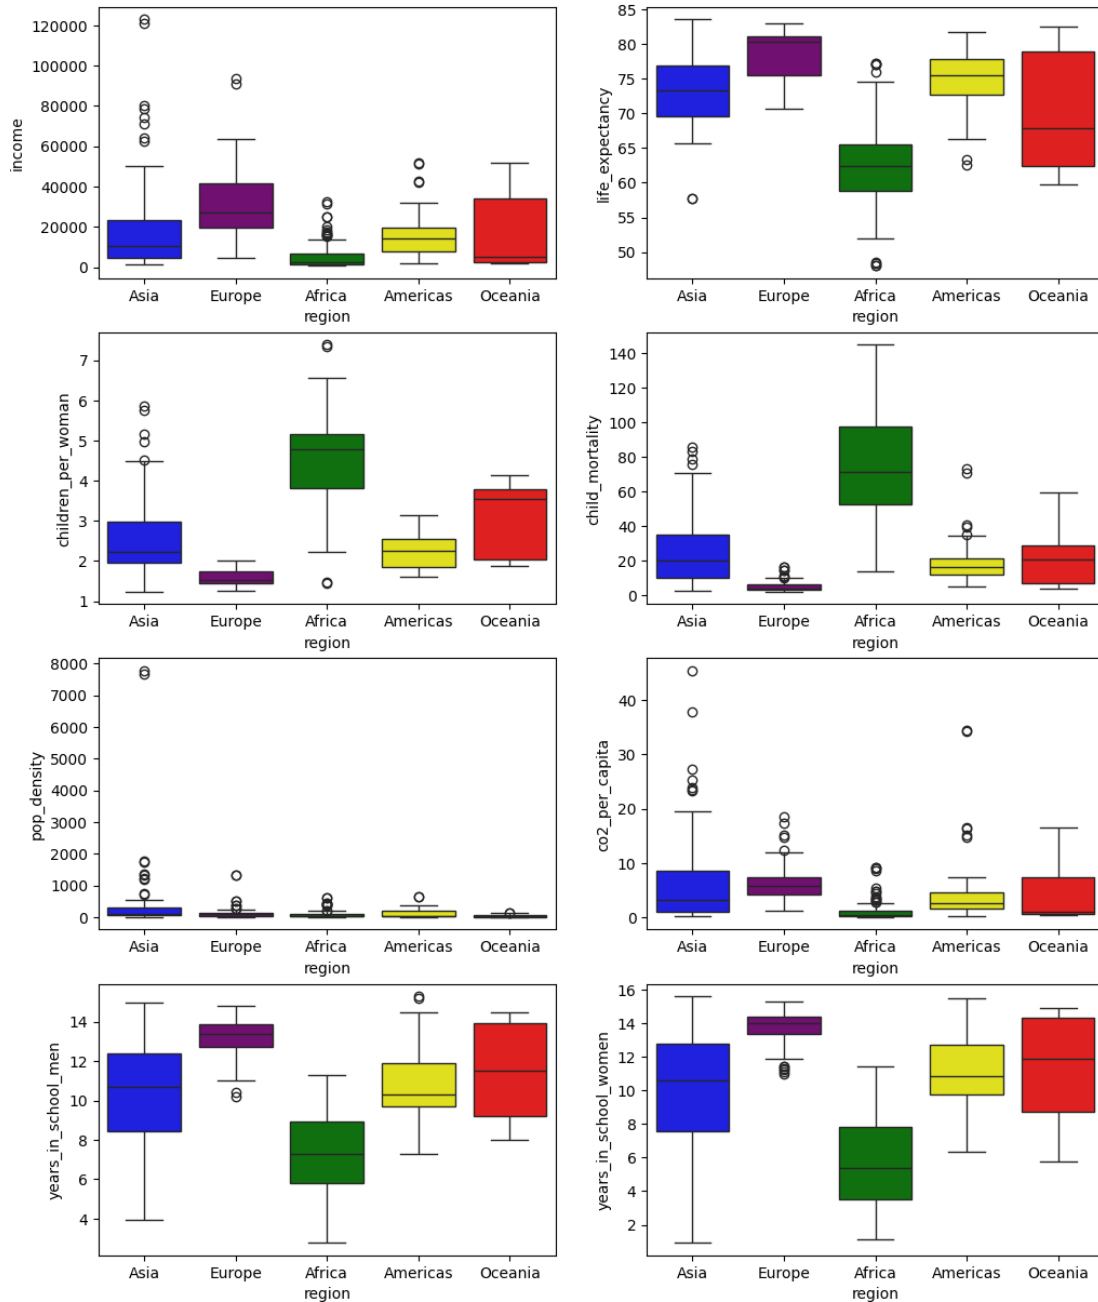
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(ax=axes[index], x="region", y=col_name, data=data_period,  
palette=region_colors)
```

[18]: Text(0.5, 0.98, 'Box Plots of Numerical Features by Region')



## Box Plots of Numerical Features by Region



Actually, these box plots show more outliers than the previous graphs. As studied in this course, the presence of outliers can be due to errors in the measurement, but it can also be right. Let's have a look at data belonging to Oceania to understand the reason behind the outliers:

```
[19]: data_period[data_period["region"] == "Oceania"]
```

```
[19]:
```

	country	year	population	region	sub_region \
358	Australia	2013	23200000	Oceania	Australia and New Zealand
359	Australia	2014	23500000	Oceania	Australia and New Zealand
2336	Fiji	2013	880000	Oceania	Melanesia
2337	Fiji	2014	886000	Oceania	Melanesia
3686	Kiribati	2013	109000	Oceania	Micronesia
3687	Kiribati	2014	110000	Oceania	Micronesia
4928	New Zealand	2013	4520000	Oceania	Australia and New Zealand
4929	New Zealand	2014	4570000	Oceania	Australia and New Zealand
5358	Papua New Guinea	2013	7590000	Oceania	Melanesia
5359	Papua New Guinea	2014	7760000	Oceania	Melanesia
5808	Samoa	2013	191000	Oceania	Polynesia
5809	Samoa	2014	192000	Oceania	Polynesia
6169	Solomon Islands	2013	564000	Oceania	Melanesia
6170	Solomon Islands	2014	576000	Oceania	Melanesia
6905	Tonga	2013	105000	Oceania	Polynesia
6906	Tonga	2014	106000	Oceania	Polynesia
7354	United States	2013	316000000	Oceania	Micronesia
7356	United States	2014	318000000	Oceania	Micronesia
7490	Vanuatu	2013	253000	Oceania	Melanesia
7491	Vanuatu	2014	259000	Oceania	Melanesia

	income_group	life_expectancy	income	children_per_woman \
358	High	82.5	42900	1.89
359	High	82.6	43400	1.87
2336	Upper middle	65.5	7980	2.59
2337	Upper middle	65.5	8350	2.57
3686	Lower middle	61.2	1830	3.77
3687	Lower middle	61.4	1840	3.73
4928	High	81.5	33800	2.05
4929	High	81.5	34500	2.03
5358	Lower middle	59.8	2470	3.81
5359	Lower middle	60.1	2620	3.76
5808	Upper middle	71.6	5490	4.15
5809	Upper middle	71.6	5510	4.09
6169	Lower middle	62.4	2030	4.03
6170	Lower middle	62.4	2020	3.97
6905	Upper middle	70.1	4950	3.77
6906	Upper middle	70.2	5030	3.72
7354	High	78.9	51000	1.96
7356	High	78.9	51800	1.95
7490	Lower middle	63.5	2890	3.38
7491	Lower middle	63.5	2890	3.35

child_mortality	pop_density	co2_per_capita	years_in_school_men \
-----------------	-------------	----------------	-----------------------

358	4.2	3.01	16.100	13.90
359	4.0	3.06	15.400	14.00
2336	23.4	48.20	1.310	11.40
2337	23.0	48.50	1.320	11.50
3686	58.8	134.00	0.574	9.36
3687	57.4	136.00	0.564	9.48
4928	5.9	17.20	7.410	14.20
4929	5.7	17.30	7.590	14.30
5358	59.5	16.80	0.815	8.01
5359	57.9	17.10	0.815	8.15
5808	18.4	67.40	1.040	12.10
5809	18.1	67.90	1.030	12.20
6169	27.2	20.10	0.358	8.63
6170	26.8	20.60	0.350	8.78
6905	17.4	146.00	1.080	11.50
6906	17.1	147.00	1.140	11.60
7354	6.9	34.50	16.400	14.40
7356	6.8	34.70	16.500	14.50
7490	29.1	20.80	0.420	9.09
7491	28.7	21.20	0.595	9.23

	years_in_school_women
358	14.30
359	14.40
2336	11.80
2337	11.90
3686	9.70
3687	9.85
4928	14.80
4929	14.90
5358	5.75
5359	5.89
5808	12.90
5809	13.00
6169	7.50
6170	7.68
6905	11.90
6906	12.00
7354	14.80
7356	14.90
7490	8.61
7491	8.78

We can see that the outliers from the income column are Australia and New Zealand. These values represent actual data and removing them would cause a loss of data. The same can be said regarding the child\_mortality, pop\_density and co2\_per\_capita besides any other feature for any other continents. Therefore, ***The outliers would not be removed.***

## 5.1 Exploratory Data Analysis

##Univariate Analysis

```
[20]: # For numerical attributes, let's describe the dataset to get mean, median, etc.  
data.describe()
```

```
[20]:
```

	year	population	life_expectancy	income \
count	7717.000000	7.717000e+03	7717.000000	7717.000000
mean	1992.435921	3.313877e+07	66.028768	12962.267850
std	12.979325	1.188602e+08	9.654891	17553.735789
min	1970.000000	5.120000e+04	12.600000	247.000000
25%	1981.000000	2.440000e+06	58.700000	2280.000000
50%	1993.000000	7.120000e+06	68.400000	6560.000000
75%	2004.000000	2.080000e+07	73.600000	16600.000000
max	2014.000000	1.390000e+09	83.600000	178000.000000


  

	children_per_woman	child_mortality	pop_density	co2_per_capita \
count	7717.000000	7717.000000	7717.000000	7717.000000
mean	3.913842	74.726785	136.687679	4.665658
std	1.990729	73.384581	417.938633	7.215037
min	1.120000	2.300000	0.823000	0.004330
25%	2.060000	17.000000	18.000000	0.421000
50%	3.480000	46.700000	53.400000	1.870000
75%	5.710000	113.000000	122.000000	6.570000
max	8.870000	399.000000	7780.000000	87.700000


	years_in_school_men	years_in_school_women
count	7717.000000	7717.000000
mean	7.720621	6.981117
std	3.190283	3.888600
min	0.900000	0.210000
25%	5.180000	3.600000
50%	7.680000	7.030000
75%	10.200000	10.100000
max	15.300000	15.600000

```
[21]: # Visualizing distributions of numerical attributes  
# Histogram for Life Expectancy  
plt.figure(figsize=(10, 6))  
sns.histplot(data['life_expectancy'], kde=True)  
plt.title('Distribution of Life Expectancy')  
plt.xlabel('Life Expectancy')  
plt.ylabel('Frequency')  
plt.show()
```



Here, we can see Life Expectancy is left-skewed, which means that most of the data points are concentrated towards the higher end of the scale, while a few extreme values (outliers) pull the distribution towards the lower end. In other words, the majority of countries or regions tend to have relatively high life expectancies, but there are a few places with significantly lower life expectancies that drag the overall distribution to the left. This skewness can occur due to various factors, such as differences in healthcare, socioeconomic conditions, and lifestyle choices etc.

```
[22]: # Histogram for Income
plt.figure(figsize=(8, 8))
sns.histplot(data["income"], kde=True) # Replace 'income' with the relevant_
    column name
plt.title("Distribution of Income")
plt.xlabel("Income")
plt.ylabel("Frequency")
plt.show()
```



The right-skewed distribution of income indicates that most data points are concentrated towards the lower end of the income scale, while a few extreme values (outliers) pull the distribution towards the higher end. In other words, the majority of individuals tend to have relatively lower incomes, but there are a few high-income outliers that stretch the overall distribution to the right. Factors contributing to this skewness include income disparities, wealth concentration, and economic inequality.

## 5.2 Bivariate Analysis

1. Line Charts
2. Scatter Plots
3. Correlation plots


```
[23]: ## Line Charts or Visualizing the trend
life_expectancy_regional_data = data.groupby(["region", "year"], as_index=False).agg({'life_expectancy': 'mean'})
print(life_expectancy_regional_data)
```

	region	year	life_expectancy
0	Africa	1970	51.029787
1	Africa	1971	51.374468
2	Africa	1972	51.364583
3	Africa	1973	52.268750
4	Africa	1974	52.650000
...	...	...	...
220	Oceania	2010	69.240000
221	Oceania	2011	69.400000
222	Oceania	2012	69.500000
223	Oceania	2013	69.700000
224	Oceania	2014	69.770000

[225 rows x 3 columns]

```
[24]: import plotly.express as px
import plotly.graph_objects as go

#Comparing Life expectancy among regions
fig = px.line(data_frame= life_expectancy_regional_data, x='year',
              y='life_expectancy', color='region', labels = {'year': 'Year',
              y='life_expectancy': 'Life Expectancy'}, title = '% Employed in the Life_
              Expectancy Across Time')
fig.show()
```




Although Europe has the highest life\_expectancy index and Africa having the lowest index, Asia seems to have the highest rate of increase in Life expectancy!

[25]: **5.2.1 Let's now take a look at our Income Plots**

```
income_region_data = data.groupby(["region", "year"], as_index=False).
    .agg({'income': 'mean'})
print(income_region_data)
```

	region	year	income
0	Africa	1970	3866.170213
1	Africa	1971	3826.659574
2	Africa	1972	3676.395833
3	Africa	1973	3712.062500
4	Africa	1974	3791.562500
..	...	...	...
220	Oceania	2010	14938.000000
221	Oceania	2011	15142.000000
222	Oceania	2012	15381.000000
223	Oceania	2013	15534.000000
224	Oceania	2014	15796.000000

```
[26]: fig = px.line(data_frame=income_region_data, x='year', y='income',
↳color='region', labels = {'year': 'Year', 'income': 'Income'}, title = '%_
↳Employed in Income Across Time')
#add annotations
fig.show()
```





### 5.2.2 Now we are able to get ideas about our initial 2 questions: How have socioeconomic indicators evolved globally over the years? Can we predict future trends in key indicators?

Which region has the highest Life Expectancy? Is it the same for Income? Europe has the highest Life Expectancy and Income value across time, and Africa at the lowest for both

### 5.3 What about Life expectancy and Income values for countries in each continent?

Let's start by grouping and aggregating our data! We will examine Life Expectancy first

```
life_exp_ctype_region_data = data.groupby(["country", "region", "year"], as_index=False).agg({"life_expectancy": "mean"})
print(life_exp_ctype_region_data)
```

	country	region	year	life_expectancy
0	Afghanistan	Asia	1970	45.8
1	Afghanistan	Asia	1971	45.9
2	Afghanistan	Asia	1972	45.9
3	Afghanistan	Asia	1973	46.0
4	Afghanistan	Asia	1974	46.1
...	...	...	...	...
7712	Zimbabwe	Africa	2010	49.6
7713	Zimbabwe	Africa	2011	51.9
7714	Zimbabwe	Africa	2012	54.1
7715	Zimbabwe	Africa	2013	55.6
7716	Zimbabwe	Africa	2014	57.0

[7717 rows x 4 columns]

Let's examine Americas' Life Expectancy


```
[28]: ame_lif_exp = life_exp_ctype_region_data[life_exp_ctype_region_data.
        region=="Americas"]
```

```
print(ame_lif_exp)
```

	country	region	year	life_expectancy
180	Antigua and Barbuda	Americas	1970	66.8
181	Antigua and Barbuda	Americas	1971	67.2
182	Antigua and Barbuda	Americas	1972	67.6
183	Antigua and Barbuda	Americas	1973	68.0
184	Antigua and Barbuda	Americas	1974	68.3
...	...	...	...	...
7532	Venezuela	Americas	2010	75.4
7533	Venezuela	Americas	2011	75.4
7534	Venezuela	Americas	2012	75.3
7535	Venezuela	Americas	2013	75.4
7536	Venezuela	Americas	2014	75.5

[1395 rows x 4 columns]

Life Expectancy in Americas (PLZ USE AUTOSCALE FOR THE LINE TO SEE)



Life Expectancy Index for United States and Canada are the highest in Americas where Haiti has the least life expectancy

### 5.3.1 Examining Income data by Country

[30]:


```
income_etry_region_data = data.groupby(['country', 'region', 'year'], as_index=False).agg({'income': 'mean'})
income_etry_region_data
```

	country	region	year	income
0	Afghanistan	Asia	1970	1180.0
1	Afghanistan	Asia	1971	1100.0

2	Afghanistan	Asia	1972	1050.0
3	Afghanistan	Asia	1973	1150.0
4	Afghanistan	Asia	1974	1180.0
...	...	...	...	...
7712	Zimbabwe	Africa	2010	1460.0
7713	Zimbabwe	Africa	2011	1660.0
7714	Zimbabwe	Africa	2012	1850.0
7715	Zimbabwe	Africa	2013	1900.0
7716	Zimbabwe	Africa	2014	1910.0

[7717 rows x 4 columns]

## Taking a look at America Again




We see that, income seems very much aligned with life expectancy. As United states and Canada again score highest in Americans regions. On the other hand, Guatemala, Haiti & Honduras' income data had not changed likely life expectancy. Haiti's income value as well as Life Expectancy are lower than Guatemala's & Honduras

### 5.3.2 Let's plot Life Expectancy against Income to have a better understanding!

*## 2 Scatter plots to see the picture*

```
fig = px.scatter(data,x = "life_expectancy", y = "income", title = 'Life_
↳expectancy and income Plot Across Time in Americas (PLZ USE AUTOSCALE FOR_
↳THE LINE TO SEE)',
                color = "region",animation_frame="year", range_x=[0,1],_
↳range_y= [0,100])
fig.show()
```



We can see higher income is highly correlated with Life expectancy and gradually increase upto a certain level.

*## Lets see the relationship among different variables*


*# Pairplot to visualize relationships*

```
sns.pairplot(data)
```

```
[33]: # Adjust the plot size
plt.figure(figsize=(10, 15))
```

*# Show the plot*

```
plt.show()
```



<Figure size 1000x1500 with 0 Axes>

```
[34]: 4 # 3 Correlation Analysis
      ## Let's calculate the correlation matrix
      correlation_matrix = data.corr()
      print(correlation_matrix)
```

	year	population	life_expectancy	income \
year	1.000000	0.045262	0.318198	0.118619
population	0.045262	1.000000	0.036585	-0.025842
life_expectancy	0.318198	0.036585	1.000000	0.534823
income	0.118619	-0.025842	0.534823	1.000000
children_per_woman	-0.383632	-0.111739	-0.810229	-0.379525

child_mortality	-0.383741	-0.037471	-0.893179	-0.443640
pop_density	0.066457	0.014499	0.165861	0.152944
co2_per_capita	-0.013448	0.007397	0.428303	0.809062
years_in_school_men	0.481321	0.052143	0.774083	0.448381
years_in_school_women	0.436158	0.010254	0.783480	0.462974

	children_per_woman	child_mortality	pop_density	\
year	-0.383632	-0.383741	0.066457	
population	-0.111739	-0.037471	0.014499	
life_expectancy	-0.810229	-0.893179	0.165861	
income	-0.379525	-0.443640	0.152944	
children_per_woman	1.000000	0.839869	-0.174382	
child_mortality	0.839869	1.000000	-0.144161	
pop_density	-0.174382	-0.144161	1.000000	
co2_per_capita	-0.362924	-0.401001	0.082978	
years_in_school_men	-0.835096	-0.796552	0.112296	
years_in_school_women	-0.854199	-0.804230	0.104472	


	co2_per_capita	years_in_school_men	\
year	-0.013448	0.481321	
population	0.007397	0.052143	
life_expectancy	0.428303	0.774083	
income	0.809062	0.448381	
children_per_woman	-0.362924	-0.835096	
child_mortality	-0.401001	-0.796552	
pop_density	0.082978	0.112296	
co2_per_capita	1.000000	0.387477	
years_in_school_men	0.387477	1.000000	
years_in_school_women	0.412187	0.972736	

	years_in_school_women
year	0.436158
population	0.010254
life_expectancy	0.783480
income	0.462974
children_per_woman	-0.854199
child_mortality	-0.804230
pop_density	0.104472
co2_per_capita	0.412187
years_in_school_men	0.972736
years_in_school_women	1.000000

<ipython-input-34-3d3b8c7d9b3a>:3: FutureWarning:

The default value of `numeric_only` in `DataFrame.corr` is deprecated. In a future version, it will default to `False`. Select only valid columns or specify the value of `numeric_only` to silence this warning.

```
[35]: ## Visualizing the correlation matrix
plt.figure(figsize=(8, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Socio-Economic Indicators', fontsize=15)
plt.show()
```



## 6 Feature Selection

Feature selection is a crucial step in machine learning and data analysis. It involves choosing a subset of relevant features (variables) from your dataset to build a more effective model. Another step of the feature engineering process is to create classes of the categorical data and the resulting



labels are encoded as integers (0, 1, 2)

```
[36]: from sklearn.model_selection import train_test_split, GridSearchCV, \
      ↪ cross_val_score
      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LinearRegression, LogisticRegression
      from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
      from sklearn.svm import SVR, SVC
      from sklearn.metrics import mean_squared_error, accuracy_score, \
      ↪ classification_report

      # Considering features for
      features = ["year", "population", "life_expectancy", "children_per_woman", \
      ↪ "child_mortality", "pop_density", "co2_per_capita", "years_in_school_men"]
      X = data[features]
      y_regression = data["income"] # Regression target
      y_classification = pd.qcut(data["life_expectancy"], q=3, labels=False) # \
      ↪ Classification target, dividing life expectancy into 3 classes
```

##Data Preprocessing

```
[37]: # Split the data into train and test sets
      X_train, X_test, y_train_reg, y_test_reg = train_test_split(X, y_regression, \
      ↪ test_size=0.2, random_state=42)
      X_train, X_test, y_train_cls, y_test_cls = train_test_split(X, \
      ↪ y_classification, test_size=0.2, random_state=42)

      # Feature scaling
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

##Model Selection, Training, and Evaluation

Regression Models

```
[38]: # Linear Regression
      lr = LinearRegression()
      lr.fit(X_train_scaled, y_train_reg)
      lr_score = cross_val_score(lr, X_train_scaled, y_train_reg, cv=5, \
      ↪ scoring="neg_mean_squared_error")
      print(f"Linear Regression RMSE: {np.sqrt(-lr_score.mean())}")

      # Random Forest Regressor
      rf_reg = RandomForestRegressor()
      rf_reg.fit(X_train_scaled, y_train_reg)
      rf_reg_score = cross_val_score(rf_reg, X_train_scaled, y_train_reg, cv=5, \
      ↪ scoring="neg_mean_squared_error")
```



```

print(f"Random Forest Regressor RMSE: {np.sqrt(-rf_reg_score.mean())}")

# Support Vector Regressor
svr = SVR()
svr.fit(X_train_scaled, y_train_reg)
svr_score = cross_val_score(svr, X_train_scaled, y_train_reg, cv=5,
    ↳scoring='neg_mean_squared_error')
print(f"SVR RMSE: {np.sqrt(-svr_score.mean())}")

```

Linear Regression RMSE: 9146.103687168454  
 Random Forest Regressor RMSE: 3461.2778357076218  
 SVR RMSE: 18803.063264451783

---

## Classification Models

```

[40]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score, classification_report
    from sklearn.model_selection import cross_val_predict

# Assuming X_test_scaled and y_test_cls are your test sets

# Logistic Regression Evaluation
log_reg_pred = cross_val_predict(log_reg, X_test_scaled, y_test_cls, cv=5)
print(f"Logistic Regression:")
print(f"Accuracy: {accuracy_score(y_test_cls, log_reg_pred)}")
print(f"Precision: {precision_score(y_test_cls, log_reg_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, log_reg_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, log_reg_pred, average='weighted')}")
print(classification_report(y_test_cls, log_reg_pred))

# Random Forest Classifier Evaluation
rf_cls_pred = cross_val_predict(rf_cls, X_test_scaled, y_test_cls, cv=5)
print(f"Random Forest Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, rf_cls_pred)}")
print(f"Precision: {precision_score(y_test_cls, rf_cls_pred,
    ↳average='weighted')}")
print(f"Recall: {recall_score(y_test_cls, rf_cls_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, rf_cls_pred, average='weighted')}")
print(classification_report(y_test_cls, rf_cls_pred))

# Support Vector Classifier Evaluation
svc_pred = cross_val_predict(svc, X_test_scaled, y_test_cls, cv=5)
print(f"Support Vector Classifier:")
print(f"Accuracy: {accuracy_score(y_test_cls, svc_pred)}")
print(f"Precision: {precision_score(y_test_cls, svc_pred, average='weighted')}")

```

```
print(f"Recall: {recall_score(y_test_cls, svc_pred, average='weighted')}")
print(f"F1-Score: {f1_score(y_test_cls, svc_pred, average='weighted')}")
print(classification_report(y_test_cls, svc_pred))
```

#### Logistic Regression:

Accuracy: 0.9715025906735751

Precision: 0.9715436731234735

Recall: 0.9715025906735751

F1-Score: 0.971518945936874

	precision	recall	f1-score	support
0	0.98	0.98	0.98	491
1	0.96	0.96	0.96	513
2	0.98	0.98	0.98	540
accuracy			0.97	1544
macro avg	0.97	0.97	0.97	1544
weighted avg	0.97	0.97	0.97	1544

#### Random Forest Classifier:

Accuracy: 0.9987046632124352

Precision: 0.9987096936465617

Recall: 0.9987046632124352

F1-Score: 0.9987046014954438

	precision	recall	f1-score	support
0	1.00	1.00	1.00	491
1	1.00	1.00	1.00	513
2	1.00	1.00	1.00	540
accuracy			1.00	1544
macro avg	1.00	1.00	1.00	1544
weighted avg	1.00	1.00	1.00	1544

#### Support Vector Classifier:

Accuracy: 0.9540155440414507

Precision: 0.9538671907249489

Recall: 0.9540155440414507

F1-Score: 0.9538865505825469

	precision	recall	f1-score	support
0	0.96	0.98	0.97	491
1	0.94	0.92	0.93	513
2	0.96	0.97	0.96	540
accuracy			0.95	1544
macro avg	0.95	0.95	0.95	1544

weighted avg	0.95	0.95	0.95	1544
--------------	------	------	------	------

### ##Finalizing the Most Accurate Model

After running the above code, we compare the RMSE for regression models and accuracy, precision, recall as well as f1-score for classification models.

##The model with the lowest RMSE (Random Forest Regressor) and/or highest accuracy, precision, recall as well as f1-score (Random Forest Classifier), respectively, should be considered the best model for this particular dataset and research questions.

This approach provides a comprehensive answer to the posed research questions, utilizing machine learning to explore socio-economic indicators and predict future trends.

## 7 Key Notes and insights

The initial analysis has helped to draw the following conclusions about the data: \* The null values exist for a reason and imputing them using any other value would lead to wrong conclusions about the data set. \* Similarly, the presence of outliers is not due to error measurement. Removing these values would lead to wrong results and inferences about the data. \* The pair plot and correlation heat map help the most in identifying the relationship of the different features.